fio Documentation

Release 3.13-4-g7114-dirty

a

Feb 19, 2019

Contents

Version: 3.13-4-g7114-dirty

Contents:

fio - Flexible I/O tester rev. 3.13

1.1 Overview and history

Fio was originally written to save me the hassle of writing special test case programs when I wanted to test a specific workload, either for performance reasons or to find/reproduce a bug. The process of writing such a test app can be tiresome, especially if you have to do it often. Hence I needed a tool that would be able to simulate a given I/O workload without resorting to writing a tailored test case again and again.

A test work load is difficult to define, though. There can be any number of processes or threads involved, and they can each be using their own way of generating I/O. You could have someone dirtying large amounts of memory in an memory mapped file, or maybe several threads issuing reads using asynchronous I/O. fio needed to be flexible enough to simulate both of these cases, and many more.

Fio spawns a number of threads or processes doing a particular type of I/O action as specified by the user. fio takes a number of global parameters, each inherited by the thread unless otherwise parameters given to them overriding that setting is given. The typical use of fio is to write a job file matching the I/O load one wants to simulate.

1.2 Source

Fio resides in a git repo, the canonical place is:

> git://git.kernel.dk/fio.git

When inside a corporate firewall, git:// URL sometimes does not work. If git:// does not work, use the http protocol instead:

> http://git.kernel.dk/fio.git

Snapshots are frequently generated and `fio-git-*.tar.gz` include the git meta data as well. Other tarballs are archives of official fio releases. Snapshots can download from:

> http://brick.kernel.dk/snaps/

There are also two official mirrors. Both of these are automatically synced with the main repository, when changes are pushed. If the main repo is down for some reason, either one of these is safe to use as a backup:

> git://git.kernel.org/pub/scm/linux/kernel/git/axboe/fio.git
>
> https://git.kernel.org/pub/scm/linux/kernel/git/axboe/fio.git

or

> git://github.com/axboe/fio.git
>
> https://github.com/axboe/fio.git

1.3 Mailing list

The fio project mailing list is meant for anything related to fio including general discussion, bug reporting, questions, and development. For bug reporting, see REPORTING-BUGS.

An automated mail detailing recent commits is automatically sent to the list at most daily. The list address is fio@vger.kernel.org, subscribe by sending an email to majordomo@vger.kernel.org with

> subscribe fio

in the body of the email. Archives can be found here:

> http://www.spinics.net/lists/fio/

and archives for the old list can be found here:

> http://maillist.kernel.dk/fio-devel/

1.4 Author

Fio was written by Jens Axboe <axboe@kernel.dk> to enable flexible testing of the Linux I/O subsystem and schedulers. He got tired of writing specific test applications to simulate a given workload, and found that the existing I/O benchmark/test tools out there weren't flexible enough to do what he wanted.

Jens Axboe <axboe@kernel.dk> 20060905

1.5 Binary packages

Debian: Starting with Debian "Squeeze", fio packages are part of the official Debian repository. http://packages.debian.org/search?keywords=fio .

Ubuntu: Starting with Ubuntu 10.04 LTS (aka "Lucid Lynx"), fio packages are part of the Ubuntu "universe" repository. http://packages.ubuntu.com/search?keywords=fio .

Red Hat, Fedora, CentOS & Co: Starting with Fedora 9/Extra Packages for Enterprise Linux 4, fio packages are part of the Fedora/EPEL repositories. https://apps.fedoraproject.org/packages/fio .

Mandriva: Mandriva has integrated fio into their package repository, so installing on that distro should be as easy as typing `urpmi fio`.

Arch Linux: An Arch Linux package is provided under the Community sub-repository: https://www.archlinux.org/packages/?sort=&q=fio

Solaris: Packages for Solaris are available from OpenCSW. Install their pkgutil tool (http://www.opencsw.org/get-it/pkgutil/) and then install fio via `pkgutil -i fio`.

Windows: Rebecca Cran <rebecca+fio@bluestop.org> has fio packages for Windows at https://www.bluestop.org/fio/ . The latest builds for Windows can also be grabbed from https://ci.appveyor.com/project/axboe/fio by clicking the latest x86 or x64 build, then selecting the ARTIFACTS tab.

BSDs: Packages for BSDs may be available from their binary package repositories. Look for a package "fio" using their binary package managers.

1.6 Building

Just type:

```
$ ./configure
$ make
$ make install
```

Note that GNU make is required. On BSDs it's available from devel/gmake within ports directory; on Solaris it's in the SUNWgmake package. On platforms where GNU make isn't the default, type `gmake` instead of `make`.

Configure will print the enabled options. Note that on Linux based platforms, the libaio development packages must be installed to use the libaio engine. Depending on distro, it is usually called libaio-devel or libaio-dev.

For gfio, gtk 2.18 (or newer), associated glib threads, and cairo are required to be installed. gfio isn't built automatically and can be enabled with a `--enable-gfio` option to configure.

To build fio with a cross-compiler:

```
$ make clean
$ make CROSS_COMPILE=/path/to/toolchain/prefix
```

Configure will attempt to determine the target platform automatically.

It's possible to build fio for ESX as well, use the `--esx` switch to configure.

1.6.1 Windows

On Windows, Cygwin (http://www.cygwin.com/) is required in order to build fio. To create an MSI installer package install WiX 3.8 from http://wixtoolset.org and run `dobuild.cmd` from the `os/windows` directory.

How to compile fio on 64-bit Windows:

1. Install Cygwin (http://www.cygwin.com/). Install **make** and all packages starting with **mingw64-x86_64**. Ensure **mingw64-x86_64-zlib** are installed if you wish to enable fio's log compression functionality.

2. Open the Cygwin Terminal.

3. Go to the fio directory (source files).

4. Run `make clean && make -j`.

To build fio for 32-bit Windows, ensure the -i686 versions of the previously mentioned -x86_64 packages are installed and run `./configure --build-32bit-win` before `make`. To build an fio that supports versions of Windows below Windows 7/Windows Server 2008 R2 also add `--target-win-ver=xp` to the end of the configure line that you run before doing `make`.

It's recommended that once built or installed, fio be run in a Command Prompt or other 'native' console such as console2, since there are known to be display and signal issues when running it under a Cygwin shell (see https://github.com/mintty/mintty/issues/56 and https://github.com/mintty/mintty/wiki/Tips#inputoutput-interaction-with-alien-programs for details).

1.6.2 Documentation

Fio uses Sphinx to generate documentation from the reStructuredText files. To build HTML formatted documentation run `make -C doc html` and direct your browser to `./doc/output/html/index.html`. To build manual page run `make -C doc man` and then `man doc/output/man/fio.1`. To see what other output formats are supported run `make -C doc help`.

1.7 Platforms

Fio works on (at least) Linux, Solaris, AIX, HP-UX, OSX, NetBSD, OpenBSD, Windows, FreeBSD, and DragonFly. Some features and/or options may only be available on some of the platforms, typically because those features only apply to that platform (like the solarisaio engine, or the splice engine on Linux).

Some features are not available on FreeBSD/Solaris even if they could be implemented, I'd be happy to take patches for that. An example of that is disk utility statistics and (I think) huge page support, support for that does exist in FreeBSD/Solaris.

Fio uses pthread mutexes for signalling and locking and some platforms do not support process shared pthread mutexes. As a result, on such platforms only threads are supported. This could be fixed with sysv ipc locking or other locking alternatives.

Other *BSD platforms are untested, but fio should work there almost out of the box. Since I don't do test runs or even compiles on those platforms, your mileage may vary. Sending me patches for other platforms is greatly appreciated. There's a lot of value in having the same test/benchmark tool available on all platforms.

Note that POSIX aio is not enabled by default on AIX. Messages like these:

```
Symbol resolution failed for /usr/lib/libc.a(posix_aio.o) because:
    Symbol _posix_kaio_rdwr (number 2) is not exported from dependent module /unix.
```

indicate one needs to enable POSIX aio. Run the following commands as root:

```
# lsdev -C -l posix_aio0
    posix_aio0 Defined  Posix Asynchronous I/O
# cfgmgr -l posix_aio0
# lsdev -C -l posix_aio0
    posix_aio0 Available  Posix Asynchronous I/O
```

POSIX aio should work now. To make the change permanent:

```
# chdev -l posix_aio0 -P -a autoconfig='available'
    posix_aio0 changed
```

1.8 Running fio

Running fio is normally the easiest part - you just give it the job file (or job files) as parameters:

```
$ fio [options] [jobfile] ...
```

and it will start doing what the *jobfile* tells it to do. You can give more than one job file on the command line, fio will serialize the running of those files. Internally that is the same as using the `stonewall` parameter described in the parameter section.

If the job file contains only one job, you may as well just give the parameters on the command line. The command line parameters are identical to the job parameters, with a few extra that control global parameters. For example, for the job file parameter `iodepth=2`, the mirror command line option would be `--iodepth 2` or `--iodepth=2`. You can also use the command line for giving more than one job entry. For each `--name` option that fio sees, it will start a new job with that name. Command line entries following a `--name` entry will apply to that job, until there are no more entries or a new `--name` entry is seen. This is similar to the job file options, where each option applies to the current job until a new [] job entry is seen.

fio does not need to run as root, except if the files or devices specified in the job section requires that. Some other options may also be restricted, such as memory locking, I/O scheduler switching, and decreasing the nice value.

If *jobfile* is specified as −, the job file will be read from standard input.

1.9 How fio works

The first step in getting fio to simulate a desired I/O workload, is writing a job file describing that specific setup. A job file may contain any number of threads and/or files – the typical contents of the job file is a *global* section defining shared parameters, and one or more job sections describing the jobs involved. When run, fio parses this file and sets everything up as described. If we break down a job from top to bottom, it contains the following basic parameters:

I/O type

Defines the I/O pattern issued to the file(s). We may only be reading sequentially from this file(s), or we may be writing randomly. Or even mixing reads and writes, sequentially or randomly. Should we be doing buffered I/O, or direct/raw I/O?

Block size

In how large chunks are we issuing I/O? This may be a single value, or it may describe a range of block sizes.

I/O size

How much data are we going to be reading/writing.

I/O engine

How do we issue I/O? We could be memory mapping the file, we could be using regular read/write, we could be using splice, async I/O, or even SG (SCSI generic sg).

I/O depth

If the I/O engine is async, how large a queuing depth do we want to maintain?

Target file/device

How many files are we spreading the workload over.

Threads, processes and job synchronization

How many threads or processes should we spread this workload over.

The above are the basic parameters defined for a workload, in addition there's a multitude of parameters that modify other aspects of how this job behaves.

1.10 Command line options

--debug=type
Enable verbose tracing *type* of various fio actions. May be `all` for all types or individual types separated by a

comma (e.g. `--debug=file,mem` will enable file and memory debugging). Currently, additional logging is available for:

process Dump info related to processes.

file Dump info related to file actions.

io Dump info related to I/O queuing.

mem Dump info related to memory allocations.

blktrace Dump info related to blktrace setup.

verify Dump info related to I/O verification.

all Enable all debug options.

random Dump info related to random offset generation.

parse Dump info related to option matching and parsing.

diskutil Dump info related to disk utilization updates.

job:x Dump info only related to job number x.

mutex Dump info only related to mutex up/down ops.

profile Dump info related to profile extensions.

time Dump info related to internal time keeping.

net Dump info related to networking connections.

rate Dump info related to I/O rate switching.

compress Dump info related to log compress/decompress.

? or help Show available debug options.

--parse-only
Parse options only, don't start any I/O.

--merge-blktrace-only
Merge blktraces only, don't start any I/O.

--output=filename
Write output to file *filename*.

--output-format=format
Set the reporting *format* to *normal*, *terse*, *json*, or *json+*. Multiple formats can be selected, separated by a comma. *terse* is a CSV based format. *json+* is like *json*, except it adds a full dump of the latency buckets.

--bandwidth-log
Generate aggregate bandwidth logs.

--minimal
Print statistics in a terse, semicolon-delimited format.

--append-terse
Print statistics in selected mode AND terse, semicolon-delimited format. **Deprecated**, use `--output-format` instead to select multiple formats.

--terse-version=version
Set terse *version* output format (default 3, or 2 or 4 or 5).

--version
Print version information and exit.

--help
> Print a summary of the command line options and exit.

--cpuclock-test
> Perform test and validation of internal CPU clock.

--crctest=[test]
> Test the speed of the built-in checksumming functions. If no argument is given, all of them are tested. Alternatively, a comma separated list can be passed, in which case the given ones are tested.

--cmdhelp=command
> Print help information for *command*. May be `all` for all commands.

--enghelp=[ioengine[,command]]
> List all commands defined by *ioengine*, or print help for *command* defined by *ioengine*. If no *ioengine* is given, list all available ioengines.

--showcmd=jobfile
> Convert *jobfile* to a set of command-line options.

--readonly
> Turn on safety read-only checks, preventing writes and trims. The `--readonly` option is an extra safety guard to prevent users from accidentally starting a write or trim workload when that is not desired. Fio will only modify the device under test if *rw=write/randwrite/rw/randrw/trim/randtrim/trimwrite* is given. This safety net can be used as an extra precaution.

--eta=when
> Specifies when real-time ETA estimate should be printed. *when* may be *always*, *never* or *auto*. *auto* is the default, it prints ETA when requested if the output is a TTY. *always* disregards the output type, and prints ETA when requested. *never* never prints ETA.

--eta-interval=time
> By default, fio requests client ETA status roughly every second. With this option, the interval is configurable. Fio imposes a minimum allowed time to avoid flooding the console, less than 250 msec is not supported.

--eta-newline=time
> Force a new line for every *time* period passed. When the unit is omitted, the value is interpreted in seconds.

--status-interval=time
> Force a full status dump of cumulative (from job start) values at *time* intervals. This option does *not* provide per-period measurements. So values such as bandwidth are running averages. When the time unit is omitted, *time* is interpreted in seconds. Note that using this option with `--output-format=json` will yield output that technically isn't valid json, since the output will be collated sets of valid json. It will need to be split into valid sets of json after the run.

--section=name
> Only run specified section *name* in job file. Multiple sections can be specified. The `--section` option allows one to combine related jobs into one file. E.g. one job file could define light, moderate, and heavy sections. Tell fio to run only the "heavy" section by giving `--section=heavy` command line option. One can also specify the "write" operations in one section and "verify" operation in another section. The `--section` option only applies to job sections. The reserved *global* section is always parsed and used.

--alloc-size=kb
> Set the internal smalloc pool size to *kb* in KiB. The `--alloc-size` switch allows one to use a larger pool size for smalloc. If running large jobs with randommap enabled, fio can run out of memory. Smalloc is an internal allocator for shared structures from a fixed size memory pool and can grow to 16 pools. The pool size defaults to 16MiB.
>
> NOTE: While running `.fio_smalloc.*` backing store files are visible in `/tmp`.

--warnings-fatal
> All fio parser warnings are fatal, causing fio to exit with an error.

--max-jobs=nr
> Set the maximum number of threads/processes to support to *nr*. NOTE: On Linux, it may be necessary to increase the shared-memory limit (/proc/sys/kernel/shmmax) if fio runs into errors while creating jobs.

--server=args
> Start a backend server, with *args* specifying what to listen to. See *Client/Server* section.

--daemonize=pidfile
> Background a fio server, writing the pid to the given *pidfile* file.

--client=hostname
> Instead of running the jobs locally, send and run them on the given *hostname* or set of *hostname*s. See *Client/Server* section.

--remote-config=file
> Tell fio server to load this local *file*.

--idle-prof=option
> Report CPU idleness. *option* is one of the following:

>> **calibrate** Run unit work calibration only and exit.

>> **system** Show aggregate system idleness and unit work.

>> **percpu** As **system** but also show per CPU idleness.

--inflate-log=log
> Inflate and output compressed *log*.

--trigger-file=file
> Execute trigger command when *file* exists.

--trigger-timeout=time
> Execute trigger at this *time*.

--trigger=command
> Set this *command* as local trigger.

--trigger-remote=command
> Set this *command* as remote trigger.

--aux-path=path
> Use the directory specified by *path* for generated state files instead of the current working directory.

Any parameters following the options will be assumed to be job files, unless they match a job file parameter. Multiple job files can be listed and each job file will be regarded as a separate group. Fio will stonewall execution between each group.

1.11 Job file format

As previously described, fio accepts one or more job files describing what it is supposed to do. The job file format is the classic ini file, where the names enclosed in [] brackets define the job name. You are free to use any ASCII name you want, except *global* which has special meaning. Following the job name is a sequence of zero or more parameters, one per line, that define the behavior of the job. If the first character in a line is a ';' or a '#', the entire line is discarded as a comment.

A *global* section sets defaults for the jobs described in that file. A job may override a *global* section parameter, and a job file may even have several *global* sections if so desired. A job is only affected by a *global* section residing above it.

The `--cmdhelp` option also lists all options. If used with a *command* argument, `--cmdhelp` will detail the given *command*.

See the *examples/* directory for inspiration on how to write job files. Note the copyright and license requirements currently apply to *examples/* files.

So let's look at a really simple job file that defines two processes, each randomly reading from a 128MiB file:

```
; -- start job file --
[global]
rw=randread
size=128m

[job1]

[job2]

; -- end job file --
```

As you can see, the job file sections themselves are empty as all the described parameters are shared. As no `filename` option is given, fio makes up a *filename* for each of the jobs as it sees fit. On the command line, this job would look as follows:

```
$ fio --name=global --rw=randread --size=128m --name=job1 --name=job2
```

Let's look at an example that has a number of processes writing randomly to files:

```
; -- start job file --
[random-writers]
ioengine=libaio
iodepth=4
rw=randwrite
bs=32k
direct=0
size=64m
numjobs=4
; -- end job file --
```

Here we have no *global* section, as we only have one job defined anyway. We want to use async I/O here, with a depth of 4 for each file. We also increased the buffer size used to 32KiB and define numjobs to 4 to fork 4 identical jobs. The result is 4 processes each randomly writing to their own 64MiB file. Instead of using the above job file, you could have given the parameters on the command line. For this case, you would specify:

```
$ fio --name=random-writers --ioengine=libaio --iodepth=4 --rw=randwrite --bs=32k --
↪direct=0 --size=64m --numjobs=4
```

When fio is utilized as a basis of any reasonably large test suite, it might be desirable to share a set of standardized settings across multiple job files. Instead of copy/pasting such settings, any section may pull in an external `filename.fio` file with *include filename* directive, as in the following example:

```
; -- start job file including.fio --
[global]
filename=/tmp/test
filesize=1m
```

(continues on next page)

```
include glob-include.fio

[test]
rw=randread
bs=4k
time_based=1
runtime=10
include test-include.fio
; -- end job file including.fio --
```

```
; -- start job file glob-include.fio --
thread=1
group_reporting=1
; -- end job file glob-include.fio --
```

```
; -- start job file test-include.fio --
ioengine=libaio
iodepth=4
; -- end job file test-include.fio --
```

Settings pulled into a section apply to that section only (except *global* section). Include directives may be nested in that any included file may contain further include directive(s). Include files may not contain [] sections.

1.11.1 Environment variables

Fio also supports environment variable expansion in job files. Any sub-string of the form `${VARNAME}` as part of an option value (in other words, on the right of the '='), will be expanded to the value of the environment variable called *VARNAME*. If no such environment variable is defined, or *VARNAME* is the empty string, the empty string will be substituted.

As an example, let's look at a sample fio invocation and job file:

```
$ SIZE=64m NUMJOBS=4 fio jobfile.fio
```

```
; -- start job file --
[random-writers]
rw=randwrite
size=${SIZE}
numjobs=${NUMJOBS}
; -- end job file --
```

This will expand to the following equivalent job file at runtime:

```
; -- start job file --
[random-writers]
rw=randwrite
size=64m
numjobs=4
; -- end job file --
```

Fio ships with a few example job files, you can also look there for inspiration.

1.11.2 Reserved keywords

Additionally, fio has a set of reserved keywords that will be replaced internally with the appropriate value. Those keywords are:

$pagesize

> The architecture page size of the running system.

$mb_memory

> Megabytes of total memory in the system.

$ncpus

> Number of online available CPUs.

These can be used on the command line or in the job file, and will be automatically substituted with the current system values when the job is run. Simple math is also supported on these keywords, so you can perform actions like:

```
size=8*$mb_memory
```

and get that properly expanded to 8 times the size of memory in the machine.

1.12 Job file parameters

This section describes in details each parameter associated with a job. Some parameters take an option of a given type, such as an integer or a string. Anywhere a numeric value is required, an arithmetic expression may be used, provided it is surrounded by parentheses. Supported operators are:

- addition (+)
- subtraction (-)
- multiplication (*)
- division (/)
- modulus (%)
- exponentiation (^)

For time values in expressions, units are microseconds by default. This is different than for time values not in expressions (not enclosed in parentheses). The following types are used:

1.12.1 Parameter types

str String: A sequence of alphanumeric characters.

time Integer with possible time suffix. Without a unit value is interpreted as seconds unless otherwise specified. Accepts a suffix of 'd' for days, 'h' for hours, 'm' for minutes, 's' for seconds, 'ms' (or 'msec') for milliseconds and 'us' (or 'usec') for microseconds. For example, use 10m for 10 minutes.

int Integer. A whole number value, which may contain an integer prefix and an integer suffix:

> [*integer prefix*] **number** [*integer suffix*]

> The optional *integer prefix* specifies the number's base. The default is decimal. *0x* specifies hexadecimal.

The optional *integer suffix* specifies the number's units, and includes an optional unit prefix and an optional unit. For quantities of data, the default unit is bytes. For quantities of time, the default unit is seconds unless otherwise specified.

With `kb_base`=1000, fio follows international standards for unit prefixes. To specify power-of-10 decimal values defined in the International System of Units (SI):

- *K* – means kilo (K) or 1000

- *M* – means mega (M) or 1000**2

- *G* – means giga (G) or 1000**3

- *T* – means tera (T) or 1000**4

- *P* – means peta (P) or 1000**5

To specify power-of-2 binary values defined in IEC 80000-13:

- *Ki* – means kibi (Ki) or 1024

- *Mi* – means mebi (Mi) or 1024**2

- *Gi* – means gibi (Gi) or 1024**3

- *Ti* – means tebi (Ti) or 1024**4

- *Pi* – means pebi (Pi) or 1024**5

With `kb_base`=1024 (the default), the unit prefixes are opposite from those specified in the SI and IEC 80000-13 standards to provide compatibility with old scripts. For example, 4k means 4096.

For quantities of data, an optional unit of 'B' may be included (e.g., 'kB' is the same as 'k').

The *integer suffix* is not case sensitive (e.g., m/mi mean mebi/mega, not milli). 'b' and 'B' both mean byte, not bit.

Examples with `kb_base`=1000:

- *4 KiB*: 4096, 4096b, 4096B, 4ki, 4kib, 4kiB, 4Ki, 4KiB

- *1 MiB*: 1048576, 1mi, 1024ki

- *1 MB*: 1000000, 1m, 1000k

- *1 TiB*: 1099511627776, 1ti, 1024gi, 1048576mi

- *1 TB*: 1000000000, 1t, 1000m, 1000000k

Examples with `kb_base`=1024 (default):

- *4 KiB*: 4096, 4096b, 4096B, 4k, 4kb, 4kB, 4K, 4KB

- *1 MiB*: 1048576, 1m, 1024k

- *1 MB*: 1000000, 1mi, 1000ki

- *1 TiB*: 1099511627776, 1t, 1024g, 1048576m

- *1 TB*: 1000000000, 1ti, 1000mi, 1000000ki

To specify times (units are not case sensitive):

- *D* – means days

- *H* – means hours

- *M* – means minutes

- *s* – or sec means seconds (default)

- *ms* – or *msec* means milliseconds

- *us* – or *usec* means microseconds

If the option accepts an upper and lower range, use a colon ':' or minus '-' to separate such values. See irange. If the lower value specified happens to be larger than the upper value the two values are swapped.

bool Boolean. Usually parsed as an integer, however only defined for true and false (1 and 0).

irange Integer range with suffix. Allows value range to be given, such as 1024-4096. A colon may also be used as the separator, e.g. 1k:4k. If the option allows two sets of ranges, they can be specified with a ',' or '/' delimiter: 1k-4k/8k-32k. Also see int.

float_list A list of floating point numbers, separated by a ':' character.

With the above in mind, here follows the complete list of fio job parameters.

1.12.2 Units

kb_base=int

Select the interpretation of unit prefixes in input parameters.

1000 Inputs comply with IEC 80000-13 and the International System of Units (SI). Use:

- power-of-2 values with IEC prefixes (e.g., KiB)

- power-of-10 values with SI prefixes (e.g., kB)

1024 Compatibility mode (default). To avoid breaking old scripts:

- power-of-2 values with SI prefixes

- power-of-10 values with IEC prefixes

See bs for more details on input parameters.

Outputs always use correct prefixes. Most outputs include both side-by-side, like:

```
bw=2383.3kB/s (2327.4KiB/s)
```

If only one value is reported, then kb_base selects the one to use:

1000 – SI prefixes

1024 – IEC prefixes

unit_base=int

Base unit for reporting. Allowed values are:

0 Use auto-detection (default).

8 Byte based.

1 Bit based.

1.12.3 Job description

name=str

ASCII name of the job. This may be used to override the name printed by fio for this job. Otherwise the job name is used. On the command line this parameter has the special purpose of also signaling the start of a new job.

description=str
> Text description of the job. Doesn't do anything except dump this text description when this job is run. It's not parsed.

loops=int
> Run the specified number of iterations of this job. Used to repeat the same workload a given number of times. Defaults to 1.

numjobs=int
> Create the specified number of clones of this job. Each clone of job is spawned as an independent thread or process. May be used to setup a larger number of threads/processes doing the same thing. Each thread is reported separately; to see statistics for all clones as a whole, use `group_reporting` in conjunction with `new_group`. See `--max-jobs`. Default: 1.

1.12.4 Time related parameters

runtime=time
> Tell fio to terminate processing after the specified period of time. It can be quite hard to determine for how long a specified job will run, so this parameter is handy to cap the total runtime to a given time. When the unit is omitted, the value is interpreted in seconds.

time_based
> If set, fio will run for the duration of the `runtime` specified even if the file(s) are completely read or written. It will simply loop over the same workload as many times as the `runtime` allows.

startdelay=irange(time)
> Delay the start of job for the specified amount of time. Can be a single value or a range. When given as a range, each thread will choose a value randomly from within the range. Value is in seconds if a unit is omitted.

ramp_time=time
> If set, fio will run the specified workload for this amount of time before logging any performance numbers. Useful for letting performance settle before logging results, thus minimizing the runtime required for stable results. Note that the `ramp_time` is considered lead in time for a job, thus it will increase the total runtime if a special timeout or `runtime` is specified. When the unit is omitted, the value is given in seconds.

clocksource=str
> Use the given clocksource as the base of timing. The supported options are:
>
> > **gettimeofday** `gettimeofday(2)`
> >
> > **clock_gettime** `clock_gettime(2)`
> >
> > **cpu** Internal CPU clock source
>
> cpu is the preferred clocksource if it is reliable, as it is very fast (and fio is heavy on time calls). Fio will automatically use this clocksource if it's supported and considered reliable on the system it is running on, unless another clocksource is specifically set. For x86/x86-64 CPUs, this means supporting TSC Invariant.

gtod_reduce=bool
> Enable all of the `gettimeofday(2)` reducing options (disable_clat, disable_slat, disable_bw_measurement) plus reduce precision of the timeout somewhat to really shrink the `gettimeofday(2)` call count. With this option enabled, we only do about 0.4% of the `gettimeofday(2)` calls we would have done if all time keeping was enabled.

gtod_cpu=int
> Sometimes it's cheaper to dedicate a single thread of execution to just getting the current time. Fio (and databases, for instance) are very intensive on `gettimeofday(2)` calls. With this option, you can set one CPU aside for doing nothing but logging current time to a shared memory location. Then the other threads/processes that run I/O workloads need only copy that segment, instead of entering the kernel with a `gettimeofday(2)`

call. The CPU set aside for doing these time calls will be excluded from other uses. Fio will manually clear it from the CPU mask of other jobs.

1.12.5 Target file/device

directory=str

> Prefix filenames with this directory. Used to place files in a different location than `./`. You can specify a number of directories by separating the names with a ':' character. These directories will be assigned equally distributed to job clones created by `numjobs` as long as they are using generated filenames. If specific *filename(s)* are set fio will use the first listed directory, and thereby matching the *filename* semantic (which generates a file for each clone if not specified, but lets all clones use the same file if set).
>
> See the `filename` option for information on how to escape ":" and "\" characters within the directory path itself.
>
> Note: To control the directory fio will use for internal state files use `--aux-path`.

filename=str

> Fio normally makes up a *filename* based on the job name, thread number, and file number (see `filename_format`). If you want to share files between threads in a job or several jobs with fixed file paths, specify a *filename* for each of them to override the default. If the ioengine is file based, you can specify a number of files by separating the names with a ':' colon. So if you wanted a job to open `/dev/sda` and `/dev/sdb` as the two working files, you would use `filename=/dev/sda:/dev/sdb`. This also means that whenever this option is specified, `nrfiles` is ignored. The size of regular files specified by this option will be `size` divided by number of files unless an explicit size is specified by `filesize`.
>
> Each colon and backslash in the wanted path must be escaped with a `\` character. For instance, if the path is `/dev/dsk/foo@3,0:c` then you would use `filename=/dev/dsk/foo@3,0\:c` and if the path is `F:\filename` then you would use `filename=F\:\\filename`.
>
> On Windows, disk devices are accessed as `\\.\PhysicalDrive0` for the first device, `\\.\PhysicalDrive1` for the second etc. Note: Windows and FreeBSD prevent write access to areas of the disk containing in-use data (e.g. filesystems).
>
> The filename "-" is a reserved name, meaning *stdin* or *stdout*. Which of the two depends on the read/write direction set.

filename_format=str

> If sharing multiple files between jobs, it is usually necessary to have fio generate the exact names that you want. By default, fio will name a file based on the default file format specification of `jobname.jobnumber.filenumber`. With this option, that can be customized. Fio will recognize and replace the following keywords in this string:
>
> > **$jobname** The name of the worker thread or process.
> >
> > **$jobnum** The incremental number of the worker thread or process.
> >
> > **$filenum** The incremental number of the file for that worker thread or process.
>
> To have dependent jobs share a set of files, this option can be set to have fio generate filenames that are shared between the two. For instance, if `testfiles.$filenum` is specified, file number 4 for any job will be named `testfiles.4`. The default of `$jobname.$jobnum.$filenum` will be used if no other format specifier is given.
>
> If you specify a path then the directories will be created up to the main directory for the file. So for example if you specify `filename_format=a/b/c/$jobnum` then the directories a/b/c will be created before the file setup part of the job. If you specify `directory` then the path will be relative that directory, otherwise it is treated as the absolute path.

unique_filename=bool
> To avoid collisions between networked clients, fio defaults to prefixing any generated filenames (with a directory specified) with the source of the client connecting. To disable this behavior, set this option to 0.

opendir=str
> Recursively open any files below directory *str*.

lockfile=str
> Fio defaults to not locking any files before it does I/O to them. If a file or file descriptor is shared, fio can serialize I/O to that file to make the end result consistent. This is usual for emulating real workloads that share files. The lock modes are:
>
>> **none** No locking. The default.
>>
>> **exclusive** Only one thread or process may do I/O at a time, excluding all others.
>>
>> **readwrite** Read-write locking on the file. Many readers may access the file at the same time, but writes get exclusive access.

nrfiles=int
> Number of files to use for this job. Defaults to 1. The size of files will be `size` divided by this unless explicit size is specified by `filesize`. Files are created for each thread separately, and each file will have a file number within its name by default, as explained in `filename` section.

openfiles=int
> Number of files to keep open at the same time. Defaults to the same as `nrfiles`, can be set smaller to limit the number simultaneous opens.

file_service_type=str
> Defines how fio decides which file from a job to service next. The following types are defined:
>
>> **random** Choose a file at random.
>>
>> **roundrobin** Round robin over opened files. This is the default.
>>
>> **sequential** Finish one file before moving on to the next. Multiple files can still be open depending on `openfiles`.
>>
>> **zipf** Use a *Zipf* distribution to decide what file to access.
>>
>> **pareto** Use a *Pareto* distribution to decide what file to access.
>>
>> **normal** Use a *Gaussian* (normal) distribution to decide what file to access.
>>
>> **gauss** Alias for normal.
>
> For *random*, *roundrobin*, and *sequential*, a postfix can be appended to tell fio how many I/Os to issue before switching to a new file. For example, specifying `file_service_type=random:8` would cause fio to issue 8 I/Os before selecting a new file at random. For the non-uniform distributions, a floating point postfix can be given to influence how the distribution is skewed. See `random_distribution` for a description of how that would work.

ioscheduler=str
> Attempt to switch the device hosting the file to the specified I/O scheduler before running.

create_serialize=bool
> If true, serialize the file creation for the jobs. This may be handy to avoid interleaving of data files, which may greatly depend on the filesystem used and even the number of processors in the system. Default: true.

create_fsync=bool
> `fsync(2)` the data file after creation. This is the default.

create_on_open=bool
> If true, don't pre-create files but allow the job's open() to create a file when it's time to do I/O. Default: false – pre-create all necessary files when the job starts.

create_only=bool
> If true, fio will only run the setup phase of the job. If files need to be laid out or updated on disk, only that will be done – the actual job contents are not executed. Default: false.

allow_file_create=bool
> If true, fio is permitted to create files as part of its workload. If this option is false, then fio will error out if the files it needs to use don't already exist. Default: true.

allow_mounted_write=bool
> If this isn't set, fio will abort jobs that are destructive (e.g. that write) to what appears to be a mounted device or partition. This should help catch creating inadvertently destructive tests, not realizing that the test will destroy data on the mounted file system. Note that some platforms don't allow writing against a mounted device regardless of this option. Default: false.

pre_read=bool
> If this is given, files will be pre-read into memory before starting the given I/O operation. This will also clear the `invalidate` flag, since it is pointless to pre-read and then drop the cache. This will only work for I/O engines that are seek-able, since they allow you to read the same data multiple times. Thus it will not work on non-seekable I/O engines (e.g. network, splice). Default: false.

unlink=bool
> Unlink the job files when done. Not the default, as repeated runs of that job would then waste time recreating the file set again and again. Default: false.

unlink_each_loop=bool
> Unlink job files after each iteration or loop. Default: false.

zonemode=str
> Accepted values are:
>
>> **none** The `zonerange`, `zonesize` and `zoneskip` parameters are ignored.
>>
>> **strided** I/O happens in a single zone until `zonesize` bytes have been transferred. After that number of bytes has been transferred processing of the next zone starts.
>>
>> **zbd** Zoned block device mode. I/O happens sequentially in each zone, even if random I/O has been selected. Random I/O happens across all zones instead of being restricted to a single zone. The `zoneskip` parameter is ignored. `zonerange` and `zonesize` must be identical.

zonerange=int
> Size of a single zone. See also `zonesize` and `zoneskip`.

zonesize=int
> For `zonemode` =strided, this is the number of bytes to transfer before skipping `zoneskip` bytes. If this parameter is smaller than `zonerange` then only a fraction of each zone with `zonerange` bytes will be accessed. If this parameter is larger than `zonerange` then each zone will be accessed multiple times before skipping to the next zone.
>
> For `zonemode` =zbd, this is the size of a single zone. The `zonerange` parameter is ignored in this mode.

zoneskip=int
> For `zonemode` =strided, the number of bytes to skip after `zonesize` bytes of data have been transferred. This parameter must be zero for `zonemode` =zbd.

read_beyond_wp=bool
> This parameter applies to `zonemode` =zbd only.

Zoned block devices are block devices that consist of multiple zones. Each zone has a type, e.g. conventional or sequential. A conventional zone can be written at any offset that is a multiple of the block size. Sequential zones must be written sequentially. The position at which a write must occur is called the write pointer. A zoned block device can be either drive managed, host managed or host aware. For host managed devices the host must ensure that writes happen sequentially. Fio recognizes host managed devices and serializes writes to sequential zones for these devices.

If a read occurs in a sequential zone beyond the write pointer then the zoned block device will complete the read without reading any data from the storage medium. Since such reads lead to unrealistically high bandwidth and IOPS numbers fio only reads beyond the write pointer if explicitly told to do so. Default: false.

max_open_zones=int
> When running a random write test across an entire drive many more zones will be open than in a typical application workload. Hence this command line option that allows to limit the number of open zones. The number of open zones is defined as the number of zones to which write commands are issued.

zone_reset_threshold=float
> A number between zero and one that indicates the ratio of logical blocks with data to the total number of logical blocks in the test above which zones should be reset periodically.

zone_reset_frequency=float
> A number between zero and one that indicates how often a zone reset should be issued if the zone reset threshold has been exceeded. A zone reset is submitted after each (1 / zone_reset_frequency) write requests. This and the previous parameter can be used to simulate garbage collection activity.

1.12.6 I/O type

direct=bool
> If value is true, use non-buffered I/O. This is usually O_DIRECT. Note that OpenBSD and ZFS on Solaris don't support direct I/O. On Windows the synchronous ioengines don't support direct I/O. Default: false.

atomic=bool
> If value is true, attempt to use atomic direct I/O. Atomic writes are guaranteed to be stable once acknowledged by the operating system. Only Linux supports O_ATOMIC right now.

buffered=bool
> If value is true, use buffered I/O. This is the opposite of the `direct` option. Defaults to true.

readwrite=str, **rw**=str
> Type of I/O pattern. Accepted values are:
>
>> **read** Sequential reads.
>>
>> **write** Sequential writes.
>>
>> **trim** Sequential trims (Linux block devices and SCSI character devices only).
>>
>> **randread** Random reads.
>>
>> **randwrite** Random writes.
>>
>> **randtrim** Random trims (Linux block devices and SCSI character devices only).
>>
>> **rw,readwrite** Sequential mixed reads and writes.
>>
>> **randrw** Random mixed reads and writes.
>>
>> **trimwrite** Sequential trim+write sequences. Blocks will be trimmed first, then the same blocks will be written to.

Fio defaults to read if the option is not specified. For the mixed I/O types, the default is to split them 50/50. For certain types of I/O the result may still be skewed a bit, since the speed may be different.

It is possible to specify the number of I/Os to do before getting a new offset by appending :<nr> to the end of the string given. For a random read, it would look like rw=randread:8 for passing in an offset modifier with a value of 8. If the suffix is used with a sequential I/O pattern, then the *<nr>* value specified will be **added** to the generated offset for each I/O turning sequential I/O into sequential I/O with holes. For instance, using rw=write:4k will skip 4k for every write. Also see the rw_sequencer option.

rw_sequencer=str
> If an offset modifier is given by appending a number to the rw=<str> line, then this option controls how that number modifies the I/O offset being generated. Accepted values are:
>
>> **sequential** Generate sequential offset.
>>
>> **identical** Generate the same offset.
>
> sequential is only useful for random I/O, where fio would normally generate a new random offset for every I/O. If you append e.g. 8 to randread, you would get a new random offset for every 8 I/Os. The result would be a seek for only every 8 I/Os, instead of for every I/O. Use rw=randread:8 to specify that. As sequential I/O is already sequential, setting sequential for that would not result in any differences. identical behaves in a similar fashion, except it sends the same offset 8 number of times before generating a new offset.

unified_rw_reporting=bool
> Fio normally reports statistics on a per data direction basis, meaning that reads, writes, and trims are accounted and reported separately. If this option is set fio sums the results and report them as "mixed" instead.

randrepeat=bool
> Seed the random number generator used for random I/O patterns in a predictable way so the pattern is repeatable across runs. Default: true.

allrandrepeat=bool
> Seed all random number generators in a predictable way so results are repeatable across runs. Default: false.

randseed=int
> Seed the random number generators based on this seed value, to be able to control what sequence of output is being generated. If not set, the random sequence depends on the randrepeat setting.

fallocate=str
> Whether pre-allocation is performed when laying down files. Accepted values are:
>
>> **none** Do not pre-allocate space.
>>
>> **native** Use a platform's native pre-allocation call but fall back to **none** behavior if it fails/is not implemented.
>>
>> **posix** Pre-allocate via *posix_fallocate(3)*.
>>
>> **keep** Pre-allocate via *fallocate(2)* with FALLOC_FL_KEEP_SIZE set.
>>
>> **0** Backward-compatible alias for **none**.
>>
>> **1** Backward-compatible alias for **posix**.
>
> May not be available on all supported platforms. **keep** is only available on Linux. If using ZFS on Solaris this cannot be set to **posix** because ZFS doesn't support pre-allocation. Default: **native** if any pre-allocation methods are available, **none** if not.

fadvise_hint=str
> Use *posix_fadvise(2)* or *posix_fadvise(2)* to advise the kernel on what I/O patterns are likely to be issued. Accepted values are:
>
>> **0** Backwards-compatible hint for "no hint".

1 Backwards compatible hint for "advise with fio workload type". This uses **FADV_RANDOM** for a random workload, and **FADV_SEQUENTIAL** for a sequential workload.

sequential Advise using **FADV_SEQUENTIAL**.

random Advise using **FADV_RANDOM**.

write_hint=str
Use `fcntl(2)` to advise the kernel what life time to expect from a write. Only supported on Linux, as of version 4.13. Accepted values are:

none No particular life time associated with this file.

short Data written to this file has a short life time.

medium Data written to this file has a medium life time.

long Data written to this file has a long life time.

extreme Data written to this file has a very long life time.

The values are all relative to each other, and no absolute meaning should be associated with them.

offset=int
Start I/O at the provided offset in the file, given as either a fixed size in bytes or a percentage. If a percentage is given, the generated offset will be aligned to the minimum `blocksize` or to the value of `offset_align` if provided. Data before the given offset will not be touched. This effectively caps the file size at *real_size - offset*. Can be combined with `size` to constrain the start and end range of the I/O workload. A percentage can be specified by a number between 1 and 100 followed by '%', for example, `offset=20%` to specify 20%.

offset_align=int
If set to non-zero value, the byte offset generated by a percentage `offset` is aligned upwards to this value. Defaults to 0 meaning that a percentage offset is aligned to the minimum block size.

offset_increment=int
If this is provided, then the real offset becomes *offset + offset_increment * thread_number*, where the thread number is a counter that starts at 0 and is incremented for each sub-job (i.e. when `numjobs` option is specified). This option is useful if there are several jobs which are intended to operate on a file in parallel disjoint segments, with even spacing between the starting points.

number_ios=int
Fio will normally perform I/Os until it has exhausted the size of the region set by `size`, or if it exhaust the allocated time (or hits an error condition). With this setting, the range/size can be set independently of the number of I/Os to perform. When fio reaches this number, it will exit normally and report status. Note that this does not extend the amount of I/O that will be done, it will only stop fio if this condition is met before other end-of-job criteria.

fsync=int
If writing to a file, issue an `fsync(2)` (or its equivalent) of the dirty data for every number of blocks given. For example, if you give 32 as a parameter, fio will sync the file after every 32 writes issued. If fio is using non-buffered I/O, we may not sync the file. The exception is the sg I/O engine, which synchronizes the disk cache anyway. Defaults to 0, which means fio does not periodically issue and wait for a sync to complete. Also see `end_fsync` and `fsync_on_close`.

fdatasync=int
Like `fsync` but uses `fdatasync(2)` to only sync data and not metadata blocks. In Windows, FreeBSD, and DragonFlyBSD there is no `fdatasync(2)` so this falls back to using `fsync(2)`. Defaults to 0, which means fio does not periodically issue and wait for a data-only sync to complete.

write_barrier=int
Make every *N-th* write a barrier write.

sync_file_range=str:int

> Use `sync_file_range(2)` for every *int* number of write operations. Fio will track range of writes that have happened since the last `sync_file_range(2)` call. *str* can currently be one or more of:
>
> > **wait_before** SYNC_FILE_RANGE_WAIT_BEFORE
> >
> > **write** SYNC_FILE_RANGE_WRITE
> >
> > **wait_after** SYNC_FILE_RANGE_WAIT_AFTER
>
> So if you do `sync_file_range=wait_before,write:8`, fio would use SYNC_FILE_RANGE_WAIT_BEFORE | SYNC_FILE_RANGE_WRITE for every 8 writes. Also see the `sync_file_range(2)` man page. This option is Linux specific.

overwrite=bool

> If true, writes to a file will always overwrite existing data. If the file doesn't already exist, it will be created before the write phase begins. If the file exists and is large enough for the specified write phase, nothing will be done. Default: false.

end_fsync=bool

> If true, `fsync(2)` file contents when a write stage has completed. Default: false.

fsync_on_close=bool

> If true, fio will `fsync(2)` a dirty file on close. This differs from end_fsync in that it will happen on every file close, not just at the end of the job. Default: false.

rwmixread=int

> Percentage of a mixed workload that should be reads. Default: 50.

rwmixwrite=int

> Percentage of a mixed workload that should be writes. If both rwmixread and rwmixwrite is given and the values do not add up to 100%, the latter of the two will be used to override the first. This may interfere with a given rate setting, if fio is asked to limit reads or writes to a certain rate. If that is the case, then the distribution may be skewed. Default: 50.

random_distribution=str:float[,str:float][,str:float]

> By default, fio will use a completely uniform random distribution when asked to perform random I/O. Sometimes it is useful to skew the distribution in specific ways, ensuring that some parts of the data is more hot than others. fio includes the following distribution models:
>
> > **random** Uniform random distribution
> >
> > **zipf** Zipf distribution
> >
> > **pareto** Pareto distribution
> >
> > **normal** Normal (Gaussian) distribution
> >
> > **zoned** Zoned random distribution
> >
> > **zoned_abs** Zone absolute random distribution
>
> When using a **zipf** or **pareto** distribution, an input value is also needed to define the access pattern. For **zipf**, this is the *Zipf theta*. For **pareto**, it's the *Pareto power*. Fio includes a test program, **fio-genzipf**, that can be used visualize what the given input values will yield in terms of hit rates. If you wanted to use **zipf** with a *theta* of 1.2, you would use `random_distribution=zipf:1.2` as the option. If a non-uniform model is used, fio will disable use of the random map. For the **normal** distribution, a normal (Gaussian) deviation is supplied as a value between 0 and 100.
>
> For a **zoned** distribution, fio supports specifying percentages of I/O access that should fall within what range of the file or device. For example, given a criteria of:
>
> - 60% of accesses should be to the first 10%

- 30% of accesses should be to the next 20%

- 8% of accesses should be to the next 30%

- 2% of accesses should be to the next 40%

we can define that through zoning of the random accesses. For the above example, the user would do:

```
random_distribution=zoned:60/10:30/20:8/30:2/40
```

A **zoned_abs** distribution works exactly like the **zoned**, except that it takes absolute sizes. For example, let's say you wanted to define access according to the following criteria:

- 60% of accesses should be to the first 20G

- 30% of accesses should be to the next 100G

- 10% of accesses should be to the next 500G

we can define an absolute zoning distribution with:

> random_distribution=zoned_abs=60/20G:30/100G:10/500g

For both **zoned** and **zoned_abs**, fio supports defining up to 256 separate zones.

Similarly to how `bssplit` works for setting ranges and percentages of block sizes. Like `bssplit`, it's possible to specify separate zones for reads, writes, and trims. If just one set is given, it'll apply to all of them. This goes for both **zoned zoned_abs** distributions.

percentage_random=int[,int][,int]
> For a random workload, set how big a percentage should be random. This defaults to 100%, in which case the workload is fully random. It can be set from anywhere from 0 to 100. Setting it to 0 would make the workload fully sequential. Any setting in between will result in a random mix of sequential and random I/O, at the given percentages. Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.

norandommap
> Normally fio will cover every block of the file when doing random I/O. If this option is given, fio will just get a new random offset without looking at past I/O history. This means that some blocks may not be read or written, and that some blocks may be read/written more than once. If this option is used with `verify` and multiple blocksizes (via `bsrange`), only intact blocks are verified, i.e., partially-overwritten blocks are ignored. With an async I/O engine and an I/O depth > 1, it is possible for the same block to be overwritten, which can cause verification errors. Either do not use norandommap in this case, or also use the lfsr random generator.

softrandommap=bool
> See `norandommap`. If fio runs with the random block map enabled and it fails to allocate the map, if this option is set it will continue without a random block map. As coverage will not be as complete as with random maps, this option is disabled by default.

random_generator=str
> Fio supports the following engines for generating I/O offsets for random I/O:

> > **tausworthe** Strong 2^88 cycle random number generator.

> > **lfsr** Linear feedback shift register generator.

> > **tausworthe64** Strong 64-bit 2^258 cycle random number generator.

> **tausworthe** is a strong random number generator, but it requires tracking on the side if we want to ensure that blocks are only read or written once. **lfsr** guarantees that we never generate the same offset twice, and it's also less computationally expensive. It's not a true random generator, however, though for I/O purposes it's typically good enough. **lfsr** only works with single block sizes, not with workloads that use multiple block sizes. If used

with such a workload, fio may read or write some blocks multiple times. The default value is **tausworthe**, unless the required space exceeds 2^32 blocks. If it does, then **tausworthe64** is selected automatically.

1.12.7 Block size

blocksize=int[,int][,int], **bs**=int[,int][,int]
> The block size in bytes used for I/O units. Default: 4096. A single value applies to reads, writes, and trims. Comma-separated values may be specified for reads, writes, and trims. A value not terminated in a comma applies to subsequent types.
>
> Examples:
>
> > **bs=256k** means 256k for reads, writes and trims.
> >
> > **bs=8k,32k** means 8k for reads, 32k for writes and trims.
> >
> > **bs=8k,32k,** means 8k for reads, 32k for writes, and default for trims.
> >
> > **bs=,8k** means default for reads, 8k for writes and trims.
> >
> > **bs=,8k,** means default for reads, 8k for writes, and default for trims.

blocksize_range=irange[,irange][,irange], **bsrange**=irange[,irange][,irange]
> A range of block sizes in bytes for I/O units. The issued I/O unit will always be a multiple of the minimum size, unless `blocksize_unaligned` is set.
>
> Comma-separated ranges may be specified for reads, writes, and trims as described in `blocksize`.
>
> Example: `bsrange=1k-4k,2k-8k`.

bssplit=str[,str][,str]
> Sometimes you want even finer grained control of the block sizes issued, not just an even split between them. This option allows you to weight various block sizes, so that you are able to define a specific amount of block sizes issued. The format for this option is:

```
bssplit=blocksize/percentage:blocksize/percentage
```

> for as many block sizes as needed. So if you want to define a workload that has 50% 64k blocks, 10% 4k blocks, and 40% 32k blocks, you would write:

```
bssplit=4k/10:64k/50:32k/40
```

> Ordering does not matter. If the percentage is left blank, fio will fill in the remaining values evenly. So a bssplit option like this one:

```
bssplit=4k/50:1k/:32k/
```

> would have 50% 4k ios, and 25% 1k and 32k ios. The percentages always add up to 100, if bssplit is given a range that adds up to more, it will error out.
>
> Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.
>
> If you want a workload that has 50% 2k reads and 50% 4k reads, while having 90% 4k writes and 10% 8k writes, you would specify:

```
bssplit=2k/50:4k/50,4k/90:8k/10
```

> Fio supports defining up to 64 different weights for each data direction.

blocksize_unaligned, bs_unaligned
> If set, fio will issue I/O units with any size within `blocksize_range`, not just multiples of the minimum size. This typically won't work with direct I/O, as that normally requires sector alignment.

bs_is_seq_rand=`bool`
> If this option is set, fio will use the normal read,write blocksize settings as sequential,random blocksize settings instead. Any random read or write will use the WRITE blocksize settings, and any sequential read or write will use the READ blocksize settings.

blockalign=`int[,int][,int]`, **ba**=`int[,int][,int]`
> Boundary to which fio will align random I/O units. Default: `blocksize`. Minimum alignment is typically 512b for using direct I/O, though it usually depends on the hardware block size. This option is mutually exclusive with using a random map for files, so it will turn off that option. Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.

1.12.8 Buffers and memory

zero_buffers
> Initialize buffers with all zeros. Default: fill buffers with random data.

refill_buffers
> If this option is given, fio will refill the I/O buffers on every submit. Only makes sense if `zero_buffers` isn't specified, naturally. Defaults to being unset i.e., the buffer is only filled at init time and the data in it is reused when possible but if any of `verify`, `buffer_compress_percentage` or `dedupe_percentage` are enabled then *refill_buffers* is also automatically enabled.

scramble_buffers=`bool`
> If `refill_buffers` is too costly and the target is using data deduplication, then setting this option will slightly modify the I/O buffer contents to defeat normal de-dupe attempts. This is not enough to defeat more clever block compression attempts, but it will stop naive dedupe of blocks. Default: true.

buffer_compress_percentage=`int`
> If this is set, then fio will attempt to provide I/O buffer content (on WRITEs) that compresses to the specified level. Fio does this by providing a mix of random data followed by fixed pattern data. The fixed pattern is either zeros, or the pattern specified by `buffer_pattern`. If the *buffer_pattern* option is used, it might skew the compression ratio slightly. Setting *buffer_compress_percentage* to a value other than 100 will also enable `refill_buffers` in order to reduce the likelihood that adjacent blocks are so similar that they over compress when seen together. See `buffer_compress_chunk` for how to set a finer or coarser granularity for the random/fixed data region. Defaults to unset i.e., buffer data will not adhere to any compression level.

buffer_compress_chunk=`int`
> This setting allows fio to manage how big the random/fixed data region is when using `buffer_compress_percentage`. When *buffer_compress_chunk* is set to some non-zero value smaller than the block size, fio can repeat the random/fixed region throughout the I/O buffer at the specified interval (which particularly useful when bigger block sizes are used for a job). When set to 0, fio will use a chunk size that matches the block size resulting in a single random/fixed region within the I/O buffer. Defaults to 512. When the unit is omitted, the value is interpreted in bytes.

buffer_pattern=`str`
> If set, fio will fill the I/O buffers with this pattern or with the contents of a file. If not set, the contents of I/O buffers are defined by the other options related to buffer contents. The setting can be any pattern of bytes, and can be prefixed with 0x for hex values. It may also be a string, where the string must then be wrapped with `""`. Or it may also be a filename, where the filename must be wrapped with `''` in which case the file is opened and read. Note that not all the file contents will be read if that would cause the buffers to overflow. So, for example:

```
buffer_pattern='filename'
```

or:

```
buffer_pattern="abcd"
```

or:

```
buffer_pattern=-12
```

or:

```
buffer_pattern=0xdeadface
```

Also you can combine everything together in any order:

```
buffer_pattern=0xdeadface"abcd"-12'filename'
```

dedupe_percentage=int
> If set, fio will generate this percentage of identical buffers when writing. These buffers will be naturally dedupable. The contents of the buffers depend on what other buffer compression settings have been set. It's possible to have the individual buffers either fully compressible, or not at all – this option only controls the distribution of unique buffers. Setting this option will also enable refill_buffers to prevent every buffer being identical.

invalidate=bool
> Invalidate the buffer/page cache parts of the files to be used prior to starting I/O if the platform and file type support it. Defaults to true. This will be ignored if pre_read is also specified for the same job.

sync=bool
> Use synchronous I/O for buffered writes. For the majority of I/O engines, this means using O_SYNC. Default: false.

iomem=str, **mem**=str
> Fio can use various types of memory as the I/O unit buffer. The allowed values are:

> > **malloc** Use memory from malloc(3) as the buffers. Default memory type.

> > **shm** Use shared memory as the buffers. Allocated through shmget(2).

> > **shmhuge** Same as shm, but use huge pages as backing.

> > **mmap** Use mmap(2) to allocate buffers. May either be anonymous memory, or can be file backed if a filename is given after the option. The format is *mem=mmap:/path/to/file*.

> > **mmaphuge** Use a memory mapped huge file as the buffer backing. Append filename after mmaphuge, ala *mem=mmaphuge:/hugetlbfs/file*.

> > **mmapshared** Same as mmap, but use a MMAP_SHARED mapping.

> > **cudamalloc** Use GPU memory as the buffers for GPUDirect RDMA benchmark. The ioengine must be *rdma*.

> The area allocated is a function of the maximum allowed bs size for the job, multiplied by the I/O depth given. Note that for **shmhuge** and **mmaphuge** to work, the system must have free huge pages allocated. This can normally be checked and set by reading/writing /proc/sys/vm/nr_hugepages on a Linux system. Fio assumes a huge page is 4MiB in size. So to calculate the number of huge pages you need for a given job file, add up the I/O depth of all jobs (normally one unless iodepth is used) and multiply by the maximum bs set. Then divide that number by the huge page size. You can see the size of the huge pages in /proc/meminfo. If no huge pages are allocated by having a non-zero number in *nr_hugepages*, using **mmaphuge** or **shmhuge** will fail. Also see hugepage-size.

> **mmaphuge** also needs to have hugetlbfs mounted and the file location should point there. So if it's mounted in /huge, you would use *mem=mmaphuge:/huge/somefile*.

iomem_align=int, **mem_align**=int

This indicates the memory alignment of the I/O memory buffers. Note that the given alignment is applied to the first I/O unit buffer, if using `iodepth` the alignment of the following buffers are given by the `bs` used. In other words, if using a `bs` that is a multiple of the page sized in the system, all buffers will be aligned to this value. If using a `bs` that is not page aligned, the alignment of subsequent I/O memory buffers is the sum of the `iomem_align` and `bs` used.

hugepage-size=int

Defines the size of a huge page. Must at least be equal to the system setting, see `/proc/meminfo`. Defaults to 4MiB. Should probably always be a multiple of megabytes, so using `hugepage-size=Xm` is the preferred way to set this to avoid setting a non-pow-2 bad value.

lockmem=int

Pin the specified amount of memory with `mlock(2)`. Can be used to simulate a smaller amount of memory. The amount specified is per worker.

1.12.9 I/O size

size=int

The total size of file I/O for each thread of this job. Fio will run until this many bytes has been transferred, unless runtime is limited by other options (such as `runtime`, for instance, or increased/decreased by `io_size`). Fio will divide this size between the available files determined by options such as `nrfiles`, `filename`, unless `filesize` is specified by the job. If the result of division happens to be 0, the size is set to the physical size of the given files or devices if they exist. If this option is not specified, fio will use the full size of the given files or devices. If the files do not exist, size must be given. It is also possible to give size as a percentage between 1 and 100. If `size=20%` is given, fio will use 20% of the full size of the given files or devices. Can be combined with `offset` to constrain the start and end range that I/O will be done within.

io_size=int, **io_limit**=int

Normally fio operates within the region set by `size`, which means that the `size` option sets both the region and size of I/O to be performed. Sometimes that is not what you want. With this option, it is possible to define just the amount of I/O that fio should do. For instance, if `size` is set to 20GiB and `io_size` is set to 5GiB, fio will perform I/O within the first 20GiB but exit when 5GiB have been done. The opposite is also possible – if `size` is set to 20GiB, and `io_size` is set to 40GiB, then fio will do 40GiB of I/O within the 0..20GiB region.

filesize=irange(int)

Individual file sizes. May be a range, in which case fio will select sizes for files at random within the given range and limited to `size` in total (if that is given). If not given, each created file is the same size. This option overrides `size` in terms of file size, which means this value is used as a fixed size or possible range of each file.

file_append=bool

Perform I/O after the end of the file. Normally fio will operate within the size of a file. If this option is set, then fio will append to the file instead. This has identical behavior to setting `offset` to the size of a file. This option is ignored on non-regular files.

fill_device=bool, **fill_fs**=bool

Sets size to something really large and waits for ENOSPC (no space left on device) as the terminating condition. Only makes sense with sequential write. For a read workload, the mount point will be filled first then I/O started on the result. This option doesn't make sense if operating on a raw device node, since the size of that is already known by the file system. Additionally, writing beyond end-of-device will not return ENOSPC there.

1.12.10 I/O engine

ioengine=str

Defines how the job issues I/O to the file. The following types are defined:

sync Basic *read(2)* or *write(2)* I/O. *lseek(2)* is used to position the I/O location. See `fsync` and `fdatasync` for syncing write I/Os.

psync Basic *pread(2)* or *pwrite(2)* I/O. Default on all supported operating systems except for Windows.

vsync Basic *readv(2)* or *writev(2)* I/O. Will emulate queuing by coalescing adjacent I/Os into a single submission.

pvsync Basic *preadv(2)* or *pwritev(2)* I/O.

pvsync2 Basic *preadv2(2)* or *pwritev2(2)* I/O.

libaio Linux native asynchronous I/O. Note that Linux may only support queued behavior with non-buffered I/O (set `direct=1` or `buffered=0`). This engine defines engine specific options.

posixaio POSIX asynchronous I/O using *aio_read(3)* and *aio_write(3)*.

solarisaio Solaris native asynchronous I/O.

windowsaio Windows native asynchronous I/O. Default on Windows.

mmap File is memory mapped with *mmap(2)* and data copied to/from using *memcpy(3)*.

splice *splice(2)* is used to transfer the data and *vmsplice(2)* to transfer data from user space to the kernel.

sg SCSI generic sg v3 I/O. May either be synchronous using the SG_IO ioctl, or if the target is an sg character device we use *read(2)* and *write(2)* for asynchronous I/O. Requires `filename` option to specify either block or character devices. This engine supports trim operations. The sg engine includes engine specific options.

null Doesn't transfer any data, just pretends to. This is mainly used to exercise fio itself and for debugging/testing purposes.

net Transfer over the network to given `host:port`. Depending on the `protocol` used, the `hostname`, `port`, `listen` and `filename` options are used to specify what sort of connection to make, while the `protocol` option determines which protocol will be used. This engine defines engine specific options.

netsplice Like **net**, but uses *splice(2)* and *vmsplice(2)* to map data and send/receive. This engine defines engine specific options.

cpuio Doesn't transfer any data, but burns CPU cycles according to the `cpuload` and `cpuchunks` options. Setting `cpuload=85` will cause that job to do nothing but burn 85% of the CPU. In case of SMP machines, use `numjobs=<nr_of_cpu>` to get desired CPU usage, as the cpuload only loads a single CPU at the desired rate. A job never finishes unless there is at least one non-cpuio job.

guasi The GUASI I/O engine is the Generic Userspace Asynchronous Syscall Interface approach to async I/O. See

http://www.xmailserver.org/guasi-lib.html

for more info on GUASI.

rdma The RDMA I/O engine supports both RDMA memory semantics (RDMA_WRITE/RDMA_READ) and channel semantics (Send/Recv) for the InfiniBand, RoCE and iWARP protocols. This engine defines engine specific options.

falloc I/O engine that does regular fallocate to simulate data transfer as fio ioengine.

DDIR_READ does fallocate(,mode = FALLOC_FL_KEEP_SIZE,).

DDIR_WRITE does fallocate(,mode = 0).

DDIR_TRIM does fallocate(,mode = FALLOC_FL_KEEP_SIZE|FALLOC_FL_PUNCH_HOLE).

ftruncate I/O engine that sends `ftruncate(2)` operations in response to write (DDIR_WRITE) events. Each ftruncate issued sets the file's size to the current block offset. `blocksize` is ignored.

e4defrag I/O engine that does regular EXT4_IOC_MOVE_EXT ioctls to simulate defragment activity in request to DDIR_WRITE event.

rados I/O engine supporting direct access to Ceph Reliable Autonomic Distributed Object Store (RADOS) via librados. This ioengine defines engine specific options.

rbd I/O engine supporting direct access to Ceph Rados Block Devices (RBD) via librbd without the need to use the kernel rbd driver. This ioengine defines engine specific options.

http I/O engine supporting GET/PUT requests over HTTP(S) with libcurl to a WebDAV or S3 endpoint. This ioengine defines engine specific options.

This engine only supports direct IO of iodepth=1; you need to scale this via numjobs. blocksize defines the size of the objects to be created.

TRIM is translated to object deletion.

gfapi Using GlusterFS libgfapi sync interface to direct access to GlusterFS volumes without having to go through FUSE. This ioengine defines engine specific options.

gfapi_async Using GlusterFS libgfapi async interface to direct access to GlusterFS volumes without having to go through FUSE. This ioengine defines engine specific options.

libhdfs Read and write through Hadoop (HDFS). The `filename` option is used to specify host,port of the hdfs name-node to connect. This engine interprets offsets a little differently. In HDFS, files once created cannot be modified so random writes are not possible. To imitate this the libhdfs engine expects a bunch of small files to be created over HDFS and will randomly pick a file from them based on the offset generated by fio backend (see the example job file to create such files, use `rw=write` option). Please note, it may be necessary to set environment variables to work with HDFS/libhdfs properly. Each job uses its own connection to HDFS.

mtd Read, write and erase an MTD character device (e.g., `/dev/mtd0`). Discards are treated as erases. Depending on the underlying device type, the I/O may have to go in a certain pattern, e.g., on NAND, writing sequentially to erase blocks and discarding before overwriting. The *trimwrite* mode works well for this constraint.

pmemblk Read and write using filesystem DAX to a file on a filesystem mounted with DAX on a persistent memory device through the PMDK libpmemblk library.

dev-dax Read and write using device DAX to a persistent memory device (e.g., /dev/dax0.0) through the PMDK libpmem library.

external Prefix to specify loading an external I/O engine object file. Append the engine filename, e.g. `ioengine=external:/tmp/foo.o` to load ioengine `foo.o` in `/tmp`. The path can be either absolute or relative. See `engines/skeleton_external.c` for details of writing an external I/O engine.

filecreate Simply create the files and do no I/O to them. You still need to set *filesize* so that all the accounting still occurs, but no actual I/O will be done other than creating the file.

libpmem Read and write using mmap I/O to a file on a filesystem mounted with DAX on a persistent memory device through the PMDK libpmem library.

ime_psync Synchronous read and write using DDN's Infinite Memory Engine (IME). This engine is very basic and issues calls to IME whenever an IO is queued.

ime_psyncv Synchronous read and write using DDN's Infinite Memory Engine (IME). This engine uses iovecs and will try to stack as much IOs as possible (if the IOs are "contiguous" and the IO depth is not exceeded) before issuing a call to IME.

ime_aio Asynchronous read and write using DDN's Infinite Memory Engine (IME). This engine will try to stack as much IOs as possible by creating requests for IME. FIO will then decide when to commit these requests.

1.12.11 I/O engine specific parameters

In addition, there are some parameters which are only valid when a specific `ioengine` is in use. These are used identically to normal parameters, with the caveat that when used on the command line, they must come after the `ioengine` that defines them is selected.

userspace_reap : `[libaio]`
Normally, with the libaio engine in use, fio will use the `io_getevents(2)` system call to reap newly returned events. With this flag turned on, the AIO ring will be read directly from user-space to reap events. The reaping mode is only enabled when polling for a minimum of 0 events (e.g. when `iodepth_batch_complete` =*0*).

hipri : `[pvsync2]`
Set RWF_HIPRI on I/O, indicating to the kernel that it's of higher priority than normal.

hipri_percentage : `[pvsync2]`
When hipri is set this determines the probability of a pvsync2 I/O being high priority. The default is 100%.

cpuload=int : `[cpuio]`
Attempt to use the specified percentage of CPU cycles. This is a mandatory option when using cpuio I/O engine.

cpuchunks=int : `[cpuio]`
Split the load into cycles of the given time. In microseconds.

exit_on_io_done=bool : `[cpuio]`
Detect when I/O threads are done, then exit.

namenode=str : `[libhdfs]`
The hostname or IP address of a HDFS cluster namenode to contact.

port=int
[libhdfs]

> The listening port of the HFDS cluster namenode.

[netsplice], [net]

> The TCP or UDP port to bind to or connect to. If this is used with `numjobs` to spawn multiple instances of the same job type, then this will be the starting port number since fio will use a range of ports.

[rdma]

> The port to use for RDMA-CM communication. This should be the same value on the client and the server side.

hostname=str : `[netsplice]` `[net]` `[rdma]`
The hostname or IP address to use for TCP, UDP or RDMA-CM based I/O. If the job is a TCP listener or UDP reader, the hostname is not used and must be omitted unless it is a valid UDP multicast address.

interface=str : `[netsplice]` `[net]`
The IP address of the network interface used to send or receive UDP multicast.

ttl=int : `[netsplice]` `[net]`
Time-to-live value for outgoing UDP multicast packets. Default: 1.

nodelay=bool : [netsplice] [net]
> Set TCP_NODELAY on TCP connections.

protocol=str, **proto**=str : [netsplice] [net]
> The network protocol to use. Accepted values are:

> **tcp** Transmission control protocol.

> **tcpv6** Transmission control protocol V6.

> **udp** User datagram protocol.

> **udpv6** User datagram protocol V6.

> **unix** UNIX domain socket.

> When the protocol is TCP or UDP, the port must also be given, as well as the hostname if the job is a TCP listener or UDP reader. For unix sockets, the normal `filename` option should be used and the port is invalid.

listen : [netsplice] [net]
> For TCP network connections, tell fio to listen for incoming connections rather than initiating an outgoing connection. The `hostname` must be omitted if this option is used.

pingpong : [netsplice] [net]
> Normally a network writer will just continue writing data, and a network reader will just consume packages. If `pingpong=1` is set, a writer will send its normal payload to the reader, then wait for the reader to send the same payload back. This allows fio to measure network latencies. The submission and completion latencies then measure local time spent sending or receiving, and the completion latency measures how long it took for the other end to receive and send back. For UDP multicast traffic `pingpong=1` should only be set for a single reader when multiple readers are listening to the same address.

window_size : [netsplice] [net]
> Set the desired socket buffer size for the connection.

mss : [netsplice] [net]
> Set the TCP maximum segment size (TCP_MAXSEG).

donorname=str : [e4defrag]
> File will be used as a block donor (swap extents between files).

inplace=int : [e4defrag]
> Configure donor file blocks allocation strategy:

> **0** Default. Preallocate donor's file on init.

> **1** Allocate space immediately inside defragment event, and free right after event.

clustername=str : [rbd,rados]
> Specifies the name of the Ceph cluster.

rbdname=str : [rbd]
> Specifies the name of the RBD.

pool=str : [rbd,rados]
> Specifies the name of the Ceph pool containing RBD or RADOS data.

clientname=str : [rbd,rados]
> Specifies the username (without the 'client.' prefix) used to access the Ceph cluster. If the *clustername* is specified, the *clientname* shall be the full *type.id* string. If no type. prefix is given, fio will add 'client.' by default.

busy_poll=bool : [rbd,rados]
> Poll store instead of waiting for completion. Usually this provides better throughput at cost of higher(up to 100%) CPU utilization.

skip_bad=bool : [mtd]
> Skip operations against known bad blocks.

hdfsdirectory : [libhdfs]
> libhdfs will create chunk in this HDFS directory.

chunk_size : [libhdfs]
> The size of the chunk to use for each file.

verb=str : [rdma]
> The RDMA verb to use on this side of the RDMA ioengine connection. Valid values are write, read, send and recv. These correspond to the equivalent RDMA verbs (e.g. write = rdma_write etc.). Note that this only needs to be specified on the client side of the connection. See the examples folder.

bindname=str : [rdma]
> The name to use to bind the local RDMA-CM connection to a local RDMA device. This could be a hostname or an IPv4 or IPv6 address. On the server side this will be passed into the rdma_bind_addr() function and on the client site it will be used in the rdma_resolve_add() function. This can be useful when multiple paths exist between the client and the server or in certain loopback configurations.

readfua=bool : [sg]
> With readfua option set to 1, read operations include the force unit access (fua) flag. Default is 0.

writefua=bool : [sg]
> With writefua option set to 1, write operations include the force unit access (fua) flag. Default is 0.

sg_write_mode=str : [sg]
> Specify the type of write commands to issue. This option can take three values:
>
> **write** This is the default where write opcodes are issued as usual.
>
> **verify** Issue WRITE AND VERIFY commands. The BYTCHK bit is set to 0. This directs the device to carry out a medium verification with no data comparison. The writefua option is ignored with this selection.
>
> **same** Issue WRITE SAME commands. This transfers a single block to the device and writes this same block of data to a contiguous sequence of LBAs beginning at the specified offset. fio's block size parameter specifies the amount of data written with each command. However, the amount of data actually transferred to the device is equal to the device's block (sector) size. For a device with 512 byte sectors, blocksize=8k will write 16 sectors with each command. fio will still generate 8k of data for each command but only the first 512 bytes will be used and transferred to the device. The writefua option is ignored with this selection.

http_host=str : [http]
> Hostname to connect to. For S3, this could be the bucket hostname. Default is **localhost**

http_user=str : [http]
> Username for HTTP authentication.

http_pass=str : [http]
> Password for HTTP authentication.

https=str : [http]
> Enable HTTPS instead of http. *on* enables HTTPS; *insecure* will enable HTTPS, but disable SSL peer verification (use with caution!). Default is **off**

http_mode=str : [http]
> Which HTTP access mode to use: *webdav*, *swift*, or *s3*. Default is **webdav**

http_s3_region=str : [http]
> The S3 region/zone string. Default is **us-east-1**

http_s3_key=str : [http]
> The S3 secret key.

---

http_s3_keyid=str : [http]
    The S3 key/access id.

http_swift_auth_token=str : [http]
    The Swift auth token. See the example configuration file on how to retrieve this.

http_verbose=int : [http]
    Enable verbose requests from libcurl. Useful for debugging. 1 turns on verbose logging from libcurl, 2 additionally enables HTTP IO tracing. Default is **0**

1.12.12 I/O depth

iodepth=int
    Number of I/O units to keep in flight against the file. Note that increasing *iodepth* beyond 1 will not affect synchronous ioengines (except for small degrees when `verify_async` is in use). Even async engines may impose OS restrictions causing the desired depth not to be achieved. This may happen on Linux when using libaio and not setting `direct`=1, since buffered I/O is not async on that OS. Keep an eye on the I/O depth distribution in the fio output to verify that the achieved depth is as expected. Default: 1.

iodepth_batch_submit=int, **iodepth_batch**=int
    This defines how many pieces of I/O to submit at once. It defaults to 1 which means that we submit each I/O as soon as it is available, but can be raised to submit bigger batches of I/O at the time. If it is set to 0 the `iodepth` value will be used.

iodepth_batch_complete_min=int, **iodepth_batch_complete**=int
    This defines how many pieces of I/O to retrieve at once. It defaults to 1 which means that we'll ask for a minimum of 1 I/O in the retrieval process from the kernel. The I/O retrieval will go on until we hit the limit set by `iodepth_low`. If this variable is set to 0, then fio will always check for completed events before queuing more I/O. This helps reduce I/O latency, at the cost of more retrieval system calls.

iodepth_batch_complete_max=int
    This defines maximum pieces of I/O to retrieve at once. This variable should be used along with `iodepth_batch_complete_min`=int variable, specifying the range of min and max amount of I/O which should be retrieved. By default it is equal to the `iodepth_batch_complete_min` value.

    Example #1:

    ```
    iodepth_batch_complete_min=1
    iodepth_batch_complete_max=<iodepth>
    ```

    which means that we will retrieve at least 1 I/O and up to the whole submitted queue depth. If none of I/O has been completed yet, we will wait.

    Example #2:

    ```
    iodepth_batch_complete_min=0
    iodepth_batch_complete_max=<iodepth>
    ```

    which means that we can retrieve up to the whole submitted queue depth, but if none of I/O has been completed yet, we will NOT wait and immediately exit the system call. In this example we simply do polling.

iodepth_low=int
    The low water mark indicating when to start filling the queue again. Defaults to the same as `iodepth`, meaning that fio will attempt to keep the queue full at all times. If `iodepth` is set to e.g. 16 and *iodepth_low* is set to 4, then after fio has filled the queue of 16 requests, it will let the depth drain down to 4 before starting to fill it again.

serialize_overlap=`bool`

> Serialize in-flight I/Os that might otherwise cause or suffer from data races. When two or more I/Os are submitted simultaneously, there is no guarantee that the I/Os will be processed or completed in the submitted order. Further, if two or more of those I/Os are writes, any overlapping region between them can become indeterminate/undefined on certain storage. These issues can cause verification to fail erratically when at least one of the racing I/Os is changing data and the overlapping region has a non-zero size. Setting `serialize_overlap` tells fio to avoid provoking this behavior by explicitly serializing in-flight I/Os that have a non-zero overlap. Note that setting this option can reduce both performance and the `iodepth` achieved.
>
> This option only applies to I/Os issued for a single job except when it is enabled along with `io_submit_mode`=offload`. In offload mode, fio will check for overlap among all I/Os submitted by offload jobs with :option:`serialize_overlap` enabled.
>
> Default: false.

io_submit_mode=`str`

> This option controls how fio submits the I/O to the I/O engine. The default is *inline*, which means that the fio job threads submit and reap I/O directly. If set to *offload*, the job threads will offload I/O submission to a dedicated pool of I/O threads. This requires some coordination and thus has a bit of extra overhead, especially for lower queue depth I/O where it can increase latencies. The benefit is that fio can manage submission rates independently of the device completion rates. This avoids skewed latency reporting if I/O gets backed up on the device side (the coordinated omission problem).

1.12.13 I/O rate

thinktime=`time`

> Stall the job for the specified period of time after an I/O has completed before issuing the next. May be used to simulate processing being done by an application. When the unit is omitted, the value is interpreted in microseconds. See `thinktime_blocks` and `thinktime_spin`.

thinktime_spin=`time`

> Only valid if `thinktime` is set - pretend to spend CPU time doing something with the data received, before falling back to sleeping for the rest of the period specified by `thinktime`. When the unit is omitted, the value is interpreted in microseconds.

thinktime_blocks=`int`

> Only valid if `thinktime` is set - control how many blocks to issue, before waiting `thinktime` usecs. If not set, defaults to 1 which will make fio wait `thinktime` usecs after every block. This effectively makes any queue depth setting redundant, since no more than 1 I/O will be queued before we have to complete it and do our `thinktime`. In other words, this setting effectively caps the queue depth if the latter is larger.

rate=`int[,int][,int]`

> Cap the bandwidth used by this job. The number is in bytes/sec, the normal suffix rules apply. Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.
>
> For example, using *rate=1m,500k* would limit reads to 1MiB/sec and writes to 500KiB/sec. Capping only reads or writes can be done with *rate=,500k* or *rate=500k,* where the former will only limit writes (to 500KiB/sec) and the latter will only limit reads.

rate_min=`int[,int][,int]`

> Tell fio to do whatever it can to maintain at least this bandwidth. Failing to meet this requirement will cause the job to exit. Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.

rate_iops=`int[,int][,int]`

> Cap the bandwidth to this number of IOPS. Basically the same as `rate`, just specified independently of bandwidth. If the job is given a block size range instead of a fixed value, the smallest block size is used as the metric. Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.

`rate_iops_min`=int[,int][,int]
> If fio doesn't meet this rate of I/O, it will cause the job to exit. Comma-separated values may be specified for reads, writes, and trims as described in `blocksize`.

`rate_process`=str
> This option controls how fio manages rated I/O submissions. The default is *linear*, which submits I/O in a linear fashion with fixed delays between I/Os that gets adjusted based on I/O completion rates. If this is set to *poisson*, fio will submit I/O based on a more real world random request flow, known as the Poisson process (https://en.wikipedia.org/wiki/Poisson_point_process). The lambda will be 10^6 / IOPS for the given workload.

`rate_ignore_thinktime`=bool
> By default, fio will attempt to catch up to the specified rate setting, if any kind of thinktime setting was used. If this option is set, then fio will ignore the thinktime and continue doing IO at the specified rate, instead of entering a catch-up mode after thinktime is done.

1.12.14 I/O latency

`latency_target`=time
> If set, fio will attempt to find the max performance point that the given workload will run at while maintaining a latency below this target. When the unit is omitted, the value is interpreted in microseconds. See `latency_window` and `latency_percentile`.

`latency_window`=time
> Used with `latency_target` to specify the sample window that the job is run at varying queue depths to test the performance. When the unit is omitted, the value is interpreted in microseconds.

`latency_percentile`=float
> The percentage of I/Os that must fall within the criteria specified by `latency_target` and `latency_window`. If not set, this defaults to 100.0, meaning that all I/Os must be equal or below to the value set by `latency_target`.

`max_latency`=time
> If set, fio will exit the job with an ETIMEDOUT error if it exceeds this maximum latency. When the unit is omitted, the value is interpreted in microseconds.

`rate_cycle`=int
> Average bandwidth for `rate` and `rate_min` over this number of milliseconds. Defaults to 1000.

1.12.15 I/O replay

`write_iolog`=str
> Write the issued I/O patterns to the specified file. See `read_iolog`. Specify a separate file for each job, otherwise the iologs will be interspersed and the file may be corrupt.

`read_iolog`=str
> Open an iolog with the specified filename and replay the I/O patterns it contains. This can be used to store a workload and replay it sometime later. The iolog given may also be a blktrace binary file, which allows fio to replay a workload captured by **blktrace**. See *blktrace(8)* for how to capture such logging data. For blktrace replay, the file needs to be turned into a blkparse binary data file first (`blkparse <device> -o /dev/null -d file_for_fio.bin`). You can specify a number of files by separating the names with a ':' character. See the `filename` option for information on how to escape ':' and '' characters within the file names. These files will be sequentially assigned to job clones created by `numjobs`.

`read_iolog_chunked`=bool
> Determines how iolog is read. If false(default) entire `read_iolog` will be read at once. If selected true, input from iolog will be read gradually. Useful when iolog is very large, or it is generated.

merge_blktrace_file=str
> When specified, rather than replaying the logs passed to `read_iolog`, the logs go through a merge phase which aggregates them into a single blktrace. The resulting file is then passed on as the `read_iolog` parameter. The intention here is to make the order of events consistent. This limits the influence of the scheduler compared to replaying multiple blktraces via concurrent jobs.

merge_blktrace_scalars=float_list
> This is a percentage based option that is index paired with the list of files passed to `read_iolog`. When merging is performed, scale the time of each event by the corresponding amount. For example, `--merge_blktrace_scalars="50:100"` runs the first trace in halftime and the second trace in realtime. This knob is separately tunable from `replay_time_scale` which scales the trace during runtime and does not change the output of the merge unlike this option.

merge_blktrace_iters=float_list
> This is a whole number option that is index paired with the list of files passed to `read_iolog`. When merging is performed, run each trace for the specified number of iterations. For example, `--merge_blktrace_iters="2:1"` runs the first trace for two iterations and the second trace for one iteration.

replay_no_stall=bool
> When replaying I/O with `read_iolog` the default behavior is to attempt to respect the timestamps within the log and replay them with the appropriate delay between IOPS. By setting this variable fio will not respect the timestamps and attempt to replay them as fast as possible while still respecting ordering. The result is the same I/O pattern to a given device, but different timings.

replay_time_scale=int
> When replaying I/O with `read_iolog`, fio will honor the original timing in the trace. With this option, it's possible to scale the time. It's a percentage option, if set to 50 it means run at 50% the original IO rate in the trace. If set to 200, run at twice the original IO rate. Defaults to 100.

replay_redirect=str
> While replaying I/O patterns using `read_iolog` the default behavior is to replay the IOPS onto the major/minor device that each IOP was recorded from. This is sometimes undesirable because on a different machine those major/minor numbers can map to a different device. Changing hardware on the same system can also result in a different major/minor mapping. `replay_redirect` causes all I/Os to be replayed onto the single specified device regardless of the device it was recorded from. i.e. `replay_redirect= /dev/sdc` would cause all I/O in the blktrace or iolog to be replayed onto `/dev/sdc`. This means multiple devices will be replayed onto a single device, if the trace contains multiple devices. If you want multiple devices to be replayed concurrently to multiple redirected devices you must blkparse your trace into separate traces and replay them with independent fio invocations. Unfortunately this also breaks the strict time ordering between multiple device accesses.

replay_align=int
> Force alignment of the byte offsets in a trace to this value. The value must be a power of 2.

replay_scale=int
> Scale byte offsets down by this factor when replaying traces. Should most likely use `replay_align` as well.

replay_skip=str
> Sometimes it's useful to skip certain IO types in a replay trace. This could be, for instance, eliminating the writes in the trace. Or not replaying the trims/discards, if you are redirecting to a device that doesn't support them. This option takes a comma separated list of read, write, trim, sync.

1.12.16 Threads, processes and job synchronization

thread
> Fio defaults to creating jobs by using fork, however if this option is given, fio will create jobs by using POSIX

Threads' function `pthread_create(3)` to create threads instead.

wait_for=str
> If set, the current job won't be started until all workers of the specified waitee job are done.
>
> `wait_for` operates on the job name basis, so there are a few limitations. First, the waitee must be defined prior to the waiter job (meaning no forward references). Second, if a job is being referenced as a waitee, it must have a unique name (no duplicate waitees).

nice=int
> Run the job with the given nice value. See man `nice(2)`.
>
> On Windows, values less than -15 set the process class to "High"; -1 through -15 set "Above Normal"; 1 through 15 "Below Normal"; and above 15 "Idle" priority class.

prio=int
> Set the I/O priority value of this job. Linux limits us to a positive value between 0 and 7, with 0 being the highest. See man `ionice(1)`. Refer to an appropriate manpage for other operating systems since meaning of priority may differ.

prioclass=int
> Set the I/O priority class. See man `ionice(1)`.

cpus_allowed=str
> Controls the same options as `cpumask`, but accepts a textual specification of the permitted CPUs instead and CPUs are indexed from 0. So to use CPUs 0 and 5 you would specify `cpus_allowed=0,5`. This option also allows a range of CPUs to be specified – say you wanted a binding to CPUs 0, 5, and 8 to 15, you would set `cpus_allowed=0,5,8-15`.
>
> On Windows, when `cpus_allowed` is unset only CPUs from fio's current processor group will be used and affinity settings are inherited from the system. An fio build configured to target Windows 7 makes options that set CPUs processor group aware and values will set both the processor group and a CPU from within that group. For example, on a system where processor group 0 has 40 CPUs and processor group 1 has 32 CPUs, `cpus_allowed` values between 0 and 39 will bind CPUs from processor group 0 and `cpus_allowed` values between 40 and 71 will bind CPUs from processor group 1. When using `cpus_allowed_policy=shared` all CPUs specified by a single `cpus_allowed` option must be from the same processor group. For Windows fio builds not built for Windows 7, CPUs will only be selected from (and be relative to) whatever processor group fio happens to be running in and CPUs from other processor groups cannot be used.

cpus_allowed_policy=str
> Set the policy of how fio distributes the CPUs specified by `cpus_allowed` or `cpumask`. Two policies are supported:
>
> > **shared** All jobs will share the CPU set specified.
> >
> > **split** Each job will get a unique CPU from the CPU set.
>
> **shared** is the default behavior, if the option isn't specified. If **split** is specified, then fio will will assign one cpu per job. If not enough CPUs are given for the jobs listed, then fio will roundrobin the CPUs in the set.

cpumask=int
> Set the CPU affinity of this job. The parameter given is a bit mask of allowed CPUs the job may run on. So if you want the allowed CPUs to be 1 and 5, you would pass the decimal value of (1 << 1 | 1 << 5), or 34. See man `sched_setaffinity(2)`. This may not work on all supported operating systems or kernel versions. This option doesn't work well for a higher CPU count than what you can store in an integer mask, so it can only control cpus 1-32. For boxes with larger CPU counts, use `cpus_allowed`.

numa_cpu_nodes=str
> Set this job running on specified NUMA nodes' CPUs. The arguments allow comma delimited list of cpu numbers, A-B ranges, or *all*. Note, to enable NUMA options support, fio must be built on a system with libnuma-dev(el) installed.

numa_mem_policy=str

> Set this job's memory policy and corresponding NUMA nodes. Format of the arguments:
>
> ```
> <mode>[:<nodelist>]
> ```
>
> mode is one of the following memory policies: default, prefer, bind, interleave or local. For default and local memory policies, no node needs to be specified. For prefer, only one node is allowed. For bind and interleave the nodelist may be as follows: a comma delimited list of numbers, A-B ranges, or *all*.

cgroup=str

> Add job to this control group. If it doesn't exist, it will be created. The system must have a mounted cgroup blkio mount point for this to work. If your system doesn't have it mounted, you can do so with:
>
> ```
> # mount -t cgroup -o blkio none /cgroup
> ```

cgroup_weight=int

> Set the weight of the cgroup to this value. See the documentation that comes with the kernel, allowed values are in the range of 100..1000.

cgroup_nodelete=bool

> Normally fio will delete the cgroups it has created after the job completion. To override this behavior and to leave cgroups around after the job completion, set cgroup_nodelete=1. This can be useful if one wants to inspect various cgroup files after job completion. Default: false.

flow_id=int

> The ID of the flow. If not specified, it defaults to being a global flow. See flow.

flow=int

> Weight in token-based flow control. If this value is used, then there is a 'flow counter' which is used to regulate the proportion of activity between two or more jobs. Fio attempts to keep this flow counter near zero. The flow parameter stands for how much should be added or subtracted to the flow counter on each iteration of the main I/O loop. That is, if one job has flow=8 and another job has flow=-1, then there will be a roughly 1:8 ratio in how much one runs vs the other.

flow_watermark=int

> The maximum value that the absolute value of the flow counter is allowed to reach before the job must wait for a lower value of the counter.

flow_sleep=int

> The period of time, in microseconds, to wait after the flow watermark has been exceeded before retrying operations.

stonewall, wait_for_previous

> Wait for preceding jobs in the job file to exit, before starting this one. Can be used to insert serialization points in the job file. A stone wall also implies starting a new reporting group, see group_reporting.

exitall

> By default, fio will continue running all other jobs when one job finishes but sometimes this is not the desired action. Setting exitall will instead make fio terminate all other jobs when one job finishes.

exec_prerun=str

> Before running this job, issue the command specified through *system(3)*. Output is redirected in a file called jobname.prerun.txt.

exec_postrun=str

> After the job completes, issue the command specified though *system(3)*. Output is redirected in a file called jobname.postrun.txt.

uid=int
> Instead of running as the invoking user, set the user ID to this value before the thread/process does any work.

gid=int
> Set group ID, see `uid`.

1.12.17 Verification

verify_only
> Do not perform specified workload, only verify data still matches previous invocation of this workload. This option allows one to check data multiple times at a later date without overwriting it. This option makes sense only for workloads that write data, and does not support workloads with the `time_based` option set.

do_verify=bool
> Run the verify phase after a write phase. Only valid if `verify` is set. Default: true.

verify=str
> If writing to a file, fio can verify the file contents after each iteration of the job. Each verification method also implies verification of special header, which is written to the beginning of each block. This header also includes meta information, like offset of the block, block number, timestamp when block was written, etc. `verify` can be combined with `verify_pattern` option. The allowed values are:
>
> > **md5** Use an md5 sum of the data area and store it in the header of each block.
> >
> > **crc64** Use an experimental crc64 sum of the data area and store it in the header of each block.
> >
> > **crc32c** Use a crc32c sum of the data area and store it in the header of each block. This will automatically use hardware acceleration (e.g. SSE4.2 on an x86 or CRC crypto extensions on ARM64) but will fall back to software crc32c if none is found. Generally the fastest checksum fio supports when hardware accelerated.
> >
> > **crc32c-intel** Synonym for crc32c.
> >
> > **crc32** Use a crc32 sum of the data area and store it in the header of each block.
> >
> > **crc16** Use a crc16 sum of the data area and store it in the header of each block.
> >
> > **crc7** Use a crc7 sum of the data area and store it in the header of each block.
> >
> > **xxhash** Use xxhash as the checksum function. Generally the fastest software checksum that fio supports.
> >
> > **sha512** Use sha512 as the checksum function.
> >
> > **sha256** Use sha256 as the checksum function.
> >
> > **sha1** Use optimized sha1 as the checksum function.
> >
> > **sha3-224** Use optimized sha3-224 as the checksum function.
> >
> > **sha3-256** Use optimized sha3-256 as the checksum function.
> >
> > **sha3-384** Use optimized sha3-384 as the checksum function.
> >
> > **sha3-512** Use optimized sha3-512 as the checksum function.
> >
> > **meta** This option is deprecated, since now meta information is included in generic verification header and meta verification happens by default. For detailed information see the description of the `verify` setting. This option is kept because of compatibility's sake with old configurations. Do not use it.

pattern Verify a strict pattern. Normally fio includes a header with some basic information and checksumming, but if this option is set, only the specific pattern set with `verify_pattern` is verified.

null Only pretend to verify. Useful for testing internals with `ioengine=null`, not for much else.

This option can be used for repeated burn-in tests of a system to make sure that the written data is also correctly read back. If the data direction given is a read or random read, fio will assume that it should verify a previously written file. If the data direction includes any form of write, the verify will be of the newly written data.

To avoid false verification errors, do not use the norandommap option when verifying data with async I/O engines and I/O depths > 1. Or use the norandommap and the lfsr random generator together to avoid writing to the same offset with muliple outstanding I/Os.

`verify_offset`=int
    Swap the verification header with data somewhere else in the block before writing. It is swapped back before verifying.

`verify_interval`=int
    Write the verification header at a finer granularity than the `blocksize`. It will be written for chunks the size of `verify_interval`. `blocksize` should divide this evenly.

`verify_pattern`=str
    If set, fio will fill the I/O buffers with this pattern. Fio defaults to filling with totally random bytes, but sometimes it's interesting to fill with a known pattern for I/O verification purposes. Depending on the width of the pattern, fio will fill 1/2/3/4 bytes of the buffer at the time (it can be either a decimal or a hex number). The `verify_pattern` if larger than a 32-bit quantity has to be a hex number that starts with either "0x" or "0X". Use with `verify`. Also, `verify_pattern` supports %o format, which means that for each block offset will be written and then verified back, e.g.:

```
verify_pattern=%o
```

Or use combination of everything:

```
verify_pattern=0xff%o"abcd"-12
```

`verify_fatal`=bool
    Normally fio will keep checking the entire contents before quitting on a block verification failure. If this option is set, fio will exit the job on the first observed failure. Default: false.

`verify_dump`=bool
    If set, dump the contents of both the original data block and the data block we read off disk to files. This allows later analysis to inspect just what kind of data corruption occurred. Off by default.

`verify_async`=int
    Fio will normally verify I/O inline from the submitting thread. This option takes an integer describing how many async offload threads to create for I/O verification instead, causing fio to offload the duty of verifying I/O contents to one or more separate threads. If using this offload option, even sync I/O engines can benefit from using an `iodepth` setting higher than 1, as it allows them to have I/O in flight while verifies are running. Defaults to 0 async threads, i.e. verification is not asynchronous.

`verify_async_cpus`=str
    Tell fio to set the given CPU affinity on the async I/O verification threads. See `cpus_allowed` for the format used.

`verify_backlog`=int
    Fio will normally verify the written contents of a job that utilizes verify once that job has completed. In other words, everything is written then everything is read back and verified. You may want to verify continually instead for a variety of reasons. Fio stores the meta data associated with an I/O block in memory, so for large

---

verify workloads, quite a bit of memory would be used up holding this meta data. If this option is enabled, fio will write only N blocks before verifying these blocks.

verify_backlog_batch=int
> Control how many blocks fio will verify if `verify_backlog` is set. If not set, will default to the value of `verify_backlog` (meaning the entire queue is read back and verified). If `verify_backlog_batch` is less than `verify_backlog` then not all blocks will be verified, if `verify_backlog_batch` is larger than `verify_backlog`, some blocks will be verified more than once.

verify_state_save=bool
> When a job exits during the write phase of a verify workload, save its current state. This allows fio to replay up until that point, if the verify state is loaded for the verify read phase. The format of the filename is, roughly:

```
<type>-<jobname>-<jobindex>-verify.state.
```

> <type> is "local" for a local run, "sock" for a client/server socket connection, and "ip" (192.168.0.1, for instance) for a networked client/server connection. Defaults to true.

verify_state_load=bool
> If a verify termination trigger was used, fio stores the current write state of each thread. This can be used at verification time so that fio knows how far it should verify. Without this information, fio will run a full verification pass, according to the settings in the job file used. Default false.

trim_percentage=int
> Number of verify blocks to discard/trim.

trim_verify_zero=bool
> Verify that trim/discarded blocks are returned as zeros.

trim_backlog=int
> Trim after this number of blocks are written.

trim_backlog_batch=int
> Trim this number of I/O blocks.

experimental_verify=bool
> Enable experimental verification.

1.12.18 Steady state

steadystate=str:float, **ss**=str:float
> Define the criterion and limit for assessing steady state performance. The first parameter designates the criterion whereas the second parameter sets the threshold. When the criterion falls below the threshold for the specified duration, the job will stop. For example, *iops_slope:0.1%* will direct fio to terminate the job when the least squares regression slope falls below 0.1% of the mean IOPS. If `group_reporting` is enabled this will apply to all jobs in the group. Below is the list of available steady state assessment criteria. All assessments are carried out using only data from the rolling collection window. Threshold limits can be expressed as a fixed value or as a percentage of the mean in the collection window.
>
> When using this feature, most jobs should include the `time_based` and `runtime` options or the `loops` option so that fio does not stop running after it has covered the full size of the specified file(s) or device(s).
>
> > **iops** Collect IOPS data. Stop the job if all individual IOPS measurements are within the specified limit of the mean IOPS (e.g., `iops:2` means that all individual IOPS values must be within 2 of the mean, whereas `iops:0.2%` means that all individual IOPS values must be within 0.2% of the mean IOPS to terminate the job).
> >
> > **iops_slope** Collect IOPS data and calculate the least squares regression slope. Stop the job if the slope falls below the specified limit.

bw Collect bandwidth data. Stop the job if all individual bandwidth measurements are within the specified limit of the mean bandwidth.

bw_slope Collect bandwidth data and calculate the least squares regression slope. Stop the job if the slope falls below the specified limit.

steadystate_duration=time, **ss_dur**=time
> A rolling window of this duration will be used to judge whether steady state has been reached. Data will be collected once per second. The default is 0 which disables steady state detection. When the unit is omitted, the value is interpreted in seconds.

steadystate_ramp_time=time, **ss_ramp**=time
> Allow the job to run for the specified duration before beginning data collection for checking the steady state job termination criterion. The default is 0. When the unit is omitted, the value is interpreted in seconds.

1.12.19 Measurements and reporting

per_job_logs=bool
> If set, this generates bw/clat/iops log with per file private filenames. If not set, jobs with identical names will share the log filename. Default: true.

group_reporting
> It may sometimes be interesting to display statistics for groups of jobs as a whole instead of for each individual job. This is especially true if numjobs is used; looking at individual thread/process output quickly becomes unwieldy. To see the final report per-group instead of per-job, use group_reporting. Jobs in a file will be part of the same reporting group, unless if separated by a stonewall, or by using new_group.

new_group
> Start a new reporting group. See: group_reporting. If not given, all jobs in a file will be part of the same reporting group, unless separated by a stonewall.

stats=bool
> By default, fio collects and shows final output results for all jobs that run. If this option is set to 0, then fio will ignore it in the final stat output.

write_bw_log=str
> If given, write a bandwidth log for this job. Can be used to store data of the bandwidth of the jobs in their lifetime.
>
> If no str argument is given, the default filename of jobname_type.x.log is used. Even when the argument is given, fio will still append the type of log. So if one specifies:

```
write_bw_log=foo
```

> The actual log name will be foo_bw.x.log where *x* is the index of the job (*1..N*, where *N* is the number of jobs). If per_job_logs is false, then the filename will not include the *.x* job index.
>
> The included **fio_generate_plots** script uses **gnuplot** to turn these text files into nice graphs. See *Log File Formats* for how data is structured within the file.

write_lat_log=str
> Same as write_bw_log, except this option creates I/O submission (e.g., name_slat.x.log), completion (e.g., name_clat.x.log), and total (e.g., name_lat.x.log) latency files instead. See write_bw_log for details about the filename format and *Log File Formats* for how data is structured within the files.

write_hist_log=str
> Same as write_bw_log but writes an I/O completion latency histogram file (e.g., name_hist.x.log) instead. Note that this file will be empty unless log_hist_msec has also been set. See write_bw_log for details about the filename format and *Log File Formats* for how data is structured within the file.

write_iops_log=str

> Same as `write_bw_log`, but writes an IOPS file (e.g. `name_iops.x.log`) instead. Because fio defaults to individual I/O logging, the value entry in the IOPS log will be 1 unless windowed logging (see `log_avg_msec`) has been enabled. See `write_bw_log` for details about the filename format and *Log File Formats* for how data is structured within the file.

log_avg_msec=int

> By default, fio will log an entry in the iops, latency, or bw log for every I/O that completes. When writing to the disk log, that can quickly grow to a very large size. Setting this option makes fio average the each log entry over the specified period of time, reducing the resolution of the log. See `log_max_value` as well. Defaults to 0, logging all entries. Also see *Log File Formats*.

log_hist_msec=int

> Same as `log_avg_msec`, but logs entries for completion latency histograms. Computing latency percentiles from averages of intervals using `log_avg_msec` is inaccurate. Setting this option makes fio log histogram entries over the specified period of time, reducing log sizes for high IOPS devices while retaining percentile accuracy. See `log_hist_coarseness` and `write_hist_log` as well. Defaults to 0, meaning histogram logging is disabled.

log_hist_coarseness=int

> Integer ranging from 0 to 6, defining the coarseness of the resolution of the histogram logs enabled with `log_hist_msec`. For each increment in coarseness, fio outputs half as many bins. Defaults to 0, for which histogram logs contain 1216 latency bins. See `write_hist_log` and *Log File Formats*.

log_max_value=bool

> If `log_avg_msec` is set, fio logs the average over that window. If you instead want to log the maximum value, set this option to 1. Defaults to 0, meaning that averaged values are logged.

log_offset=bool

> If this is set, the iolog options will include the byte offset for the I/O entry as well as the other data values. Defaults to 0 meaning that offsets are not present in logs. Also see *Log File Formats*.

log_compression=int

> If this is set, fio will compress the I/O logs as it goes, to keep the memory footprint lower. When a log reaches the specified size, that chunk is removed and compressed in the background. Given that I/O logs are fairly highly compressible, this yields a nice memory savings for longer runs. The downside is that the compression will consume some background CPU cycles, so it may impact the run. This, however, is also true if the logging ends up consuming most of the system memory. So pick your poison. The I/O logs are saved normally at the end of a run, by decompressing the chunks and storing them in the specified log file. This feature depends on the availability of zlib.

log_compression_cpus=str

> Define the set of CPUs that are allowed to handle online log compression for the I/O jobs. This can provide better isolation between performance sensitive jobs, and background compression work. See `cpus_allowed` for the format used.

log_store_compressed=bool

> If set, fio will store the log files in a compressed format. They can be decompressed with fio, using the `--inflate-log` command line parameter. The files will be stored with a `.fz` suffix.

log_unix_epoch=bool

> If set, fio will log Unix timestamps to the log files produced by enabling write_type_log for each log type, instead of the default zero-based timestamps.

block_error_percentiles=bool

> If set, record errors in trim block-sized units from writes and trims and output a histogram of how many trims it took to get to errors, and what kind of error was encountered.

bwavgtime=int

Average the calculated bandwidth over the given time. Value is specified in milliseconds. If the job also does bandwidth logging through `write_bw_log`, then the minimum of this option and `log_avg_msec` will be used. Default: 500ms.

iopsavgtime=int
    Average the calculated IOPS over the given time. Value is specified in milliseconds. If the job also does IOPS logging through `write_iops_log`, then the minimum of this option and `log_avg_msec` will be used. Default: 500ms.

disk_util=bool
    Generate disk utilization statistics, if the platform supports it. Default: true.

disable_lat=bool
    Disable measurements of total latency numbers. Useful only for cutting back the number of calls to *gettimeofday(2)*, as that does impact performance at really high IOPS rates. Note that to really get rid of a large amount of these calls, this option must be used with `disable_slat` and `disable_bw_measurement` as well.

disable_clat=bool
    Disable measurements of completion latency numbers. See `disable_lat`.

disable_slat=bool
    Disable measurements of submission latency numbers. See `disable_lat`.

disable_bw_measurement=bool, **disable_bw**=bool
    Disable measurements of throughput/bandwidth numbers. See `disable_lat`.

clat_percentiles=bool
    Enable the reporting of percentiles of completion latencies. This option is mutually exclusive with `lat_percentiles`.

lat_percentiles=bool
    Enable the reporting of percentiles of I/O latencies. This is similar to `clat_percentiles`, except that this includes the submission latency. This option is mutually exclusive with `clat_percentiles`.

percentile_list=float_list
    Overwrite the default list of percentiles for completion latencies and the block error histogram. Each number is a floating number in the range (0,100], and the maximum length of the list is 20. Use `:` to separate the numbers, and list the numbers in ascending order. For example, `--percentile_list=99.5:99.9` will cause fio to report the values of completion latency below which 99.5% and 99.9% of the observed latencies fell, respectively.

significant_figures=int
    If using `--output-format` of *normal*, set the significant figures to this value. Higher values will yield more precise IOPS and throughput units, while lower values will round. Requires a minimum value of 1 and a maximum value of 10. Defaults to 4.

1.12.20 Error handling

exitall_on_error
    When one job finishes in error, terminate the rest. The default is to wait for each job to finish.

continue_on_error=str
    Normally fio will exit the job on the first observed failure. If this option is set, fio will continue the job when there is a 'non-fatal error' (EIO or EILSEQ) until the runtime is exceeded or the I/O size specified is completed. If this option is used, there are two more stats that are appended, the total error count and the first error. The error field given in the stats is the first error that was hit during the run.

    The allowed values are:

       none Exit on any I/O or verify errors.

       read Continue on read errors, exit on all others.

       write Continue on write errors, exit on all others.

       io Continue on any I/O error, exit on all others.

       verify Continue on verify errors, exit on all others.

       all Continue on all errors.

       0 Backward-compatible alias for 'none'.

       1 Backward-compatible alias for 'all'.

ignore_error=str
> Sometimes you want to ignore some errors during test in that case you can specify error list for each error type, instead of only being able to ignore the default 'non-fatal error' using `continue_on_error`. `ignore_error=READ_ERR_LIST,WRITE_ERR_LIST,VERIFY_ERR_LIST` errors for given error type is separated with ':'. Error may be symbol ('ENOSPC', 'ENOMEM') or integer. Example:

```
ignore_error=EAGAIN,ENOSPC:122
```

> This option will ignore EAGAIN from READ, and ENOSPC and 122(EDQUOT) from WRITE. This option works by overriding `continue_on_error` with the list of errors for each error type if any.

error_dump=bool
> If set dump every error even if it is non fatal, true by default. If disabled only fatal error will be dumped.

1.13 Running predefined workloads

Fio includes predefined profiles that mimic the I/O workloads generated by other tools.

profile=str
> The predefined workload to run. Current profiles are:

       tiobench Threaded I/O bench (tiotest/tiobench) like workload.

       act Aerospike Certification Tool (ACT) like workload.

To view a profile's additional options use `--cmdhelp` after specifying the profile. For example:

```
$ fio --profile=act --cmdhelp
```

1.13.1 Act profile options

device-names=str
> Devices to use.

load=int
> ACT load multiplier. Default: 1.

test-duration=time
> How long the entire test takes to run. When the unit is omitted, the value is given in seconds. Default: 24h.

threads-per-queue=int
> Number of read I/O threads per device. Default: 8.

read-req-num-512-blocks=int
>   Number of 512B blocks to read at the time. Default: 3.

large-block-op-kbytes=int
>   Size of large block ops in KiB (writes). Default: 131072.

prep
>   Set to run ACT prep phase.

1.13.2 Tiobench profile options

size=str
>   Size in MiB.

block=int
>   Block size in bytes. Default: 4096.

numruns=int
>   Number of runs.

dir=str
>   Test directory.

threads=int
>   Number of threads.

1.14 Interpreting the output

Fio spits out a lot of output. While running, fio will display the status of the jobs created. An example of that would be:

```
Jobs: 1 (f=1): [_(1),M(1)][24.8%][r=20.5MiB/s,w=23.5MiB/s][r=82,w=94 IOPS][eta
→01m:31s]
```

The characters inside the first set of square brackets denote the current status of each thread. The first character is the first job defined in the job file, and so forth. The possible values (in typical life cycle order) are:

Idle	Run	
P		Thread setup, but not started.
C		Thread created.
I		Thread initialized, waiting or generating necessary data.
	p	Thread running pre-reading file(s).
	/	Thread is in ramp period.
	R	Running, doing sequential reads.
	r	Running, doing random reads.
	W	Running, doing sequential writes.
	w	Running, doing random writes.
	M	Running, doing mixed sequential reads/writes.
	m	Running, doing mixed random reads/writes.
	D	Running, doing sequential trims.
	d	Running, doing random trims.
	F	Running, currently waiting for *fsync(2)*.
	V	Running, doing verification of written data.
f		Thread finishing.
E		Thread exited, not reaped by main thread yet.
_		Thread reaped.
X		Thread reaped, exited with an error.
K		Thread reaped, exited due to signal.

Fio will condense the thread string as not to take up more space on the command line than needed. For instance, if you have 10 readers and 10 writers running, the output would look like this:

```
Jobs: 20 (f=20): [R(10),W(10)][4.0%][r=20.5MiB/s,w=23.5MiB/s][r=82,w=94 IOPS][eta␣
→57m:36s]
```

Note that the status string is displayed in order, so it's possible to tell which of the jobs are currently doing what. In the example above this means that jobs 1–10 are readers and 11–20 are writers.

The other values are fairly self explanatory – number of threads currently running and doing I/O, the number of currently open files (f=), the estimated completion percentage, the rate of I/O since last check (read speed listed first, then write speed and optionally trim speed) in terms of bandwidth and IOPS, and time to completion for the current running group. It's impossible to estimate runtime of the following groups (if any).

When fio is done (or interrupted by `Ctrl-C`), it will show the data for each thread, group of threads, and disks in that order. For each overall thread (or group) the output looks like:

```
Client1: (groupid=0, jobs=1): err= 0: pid=16109: Sat Jun 24 12:07:54 2017
  write: IOPS=88, BW=623KiB/s (638kB/s)(30.4MiB/50032msec)
    slat (nsec): min=500, max=145500, avg=8318.00, stdev=4781.50
    clat (usec): min=170, max=78367, avg=4019.02, stdev=8293.31
     lat (usec): min=174, max=78375, avg=4027.34, stdev=8291.79
    clat percentiles (usec):
     |  1.00th=[  302],  5.00th=[  326], 10.00th=[  343], 20.00th=[  363],
     | 30.00th=[  392], 40.00th=[  404], 50.00th=[  416], 60.00th=[  445],
     | 70.00th=[  816], 80.00th=[ 6718], 90.00th=[12911], 95.00th=[21627],
     | 99.00th=[43779], 99.50th=[51643], 99.90th=[68682], 99.95th=[72877],
     | 99.99th=[78119]
  bw (  KiB/s): min=  532, max=  686, per=0.10%, avg=622.87, stdev=24.82, samples= ␣
→100
   iops        : min=   76, max=   98, avg=88.98, stdev= 3.54, samples=  100
  lat (usec)   : 250=0.04%, 500=64.11%, 750=4.81%, 1000=2.79%
```

(continues on next page)

```
lat (msec)    : 2=4.16%, 4=1.84%, 10=4.90%, 20=11.33%, 50=5.37%
lat (msec)    : 100=0.65%
cpu           : usr=0.27%, sys=0.18%, ctx=12072, majf=0, minf=21
IO depths     : 1=85.0%, 2=13.1%, 4=1.8%, 8=0.1%, 16=0.0%, 32=0.0%, >=64=0.0%
   submit     : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
   complete   : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
   issued rwt: total=0,4450,0, short=0,0,0, dropped=0,0,0
   latency    : target=0, window=0, percentile=100.00%, depth=8
```

The job name (or first job's name when using `group_reporting`) is printed, along with the group id, count of jobs being aggregated, last error id seen (which is 0 when there are no errors), pid/tid of that thread and the time the job/group completed. Below are the I/O statistics for each data direction performed (showing writes in the example above). In the order listed, they denote:

read/write/trim The string before the colon shows the I/O direction the statistics are for. **IOPS** is the average I/Os performed per second. **BW** is the average bandwidth rate shown as: value in power of 2 format (value in power of 10 format). The last two values show: (**total I/O performed** in power of 2 format / **runtime** of that thread).

slat Submission latency (**min** being the minimum, **max** being the maximum, **avg** being the average, **stdev** being the standard deviation). This is the time it took to submit the I/O. For sync I/O this row is not displayed as the slat is really the completion latency (since queue/complete is one operation there). This value can be in nanoseconds, microseconds or milliseconds — fio will choose the most appropriate base and print that (in the example above nanoseconds was the best scale). Note: in `--minimal` mode latencies are always expressed in microseconds.

clat Completion latency. Same names as slat, this denotes the time from submission to completion of the I/O pieces. For sync I/O, clat will usually be equal (or very close) to 0, as the time from submit to complete is basically just CPU time (I/O has already been done, see slat explanation).

lat Total latency. Same names as slat and clat, this denotes the time from when fio created the I/O unit to completion of the I/O operation.

bw Bandwidth statistics based on samples. Same names as the xlat stats, but also includes the number of samples taken (**samples**) and an approximate percentage of total aggregate bandwidth this thread received in its group (**per**). This last value is only really useful if the threads in this group are on the same disk, since they are then competing for disk access.

iops IOPS statistics based on samples. Same names as bw.

lat (nsec/usec/msec) The distribution of I/O completion latencies. This is the time from when I/O leaves fio and when it gets completed. Unlike the separate read/write/trim sections above, the data here and in the remaining sections apply to all I/Os for the reporting group. 250=0.04% means that 0.04% of the I/Os completed in under 250us. 500=64.11% means that 64.11% of the I/Os required 250 to 499us for completion.

cpu CPU usage. User and system time, along with the number of context switches this thread went through, usage of system and user time, and finally the number of major and minor page faults. The CPU utilization numbers are averages for the jobs in that reporting group, while the context and fault counters are summed.

IO depths The distribution of I/O depths over the job lifetime. The numbers are divided into powers of 2 and each entry covers depths from that value up to those that are lower than the next entry – e.g., 16= covers depths from 16 to 31. Note that the range covered by a depth distribution entry can be different to the range covered by the equivalent submit/complete distribution entry.

IO submit How many pieces of I/O were submitting in a single submit call. Each entry denotes that amount and below, until the previous entry – e.g., 16=100% means that we submitted anywhere between 9 to 16 I/Os per submit call. Note that the range covered by a submit distribution entry can be different to the range covered by the equivalent depth distribution entry.

IO complete Like the above submit number, but for completions instead.

IO issued rwt The number of read/write/trim requests issued, and how many of them were short or dropped.

IO latency These values are for `latency_target` and related options. When these options are engaged, this section describes the I/O depth required to meet the specified latency target.

After each client has been listed, the group statistics are printed. They will look like this:

```
Run status group 0 (all jobs):
   READ: bw=20.9MiB/s (21.9MB/s), 10.4MiB/s-10.8MiB/s (10.9MB/s-11.3MB/s), io=64.0MiB␣
→(67.1MB), run=2973-3069msec
  WRITE: bw=1231KiB/s (1261kB/s), 616KiB/s-621KiB/s (630kB/s-636kB/s), io=64.0MiB (67.
→1MB), run=52747-53223msec
```

For each data direction it prints:

bw Aggregate bandwidth of threads in this group followed by the minimum and maximum bandwidth of all the threads in this group. Values outside of brackets are power-of-2 format and those within are the equivalent value in a power-of-10 format.

io Aggregate I/O performed of all threads in this group. The format is the same as bw.

run The smallest and longest runtimes of the threads in this group.

And finally, the disk statistics are printed. This is Linux specific. They will look like this:

```
Disk stats (read/write):
  sda: ios=16398/16511, merge=30/162, ticks=6853/819634, in_queue=826487, util=100.00%
```

Each value is printed for both reads and writes, with reads first. The numbers denote:

ios Number of I/Os performed by all groups.

merge Number of merges performed by the I/O scheduler.

ticks Number of ticks we kept the disk busy.

in_queue Total time spent in the disk queue.

util The disk utilization. A value of 100% means we kept the disk busy constantly, 50% would be a disk idling half of the time.

It is also possible to get fio to dump the current output while it is running, without terminating the job. To do that, send fio the **USR1** signal. You can also get regularly timed dumps by using the `--status-interval` parameter, or by creating a file in `/tmp` named `fio-dump-status`. If fio sees this file, it will unlink it and dump the current output status.

1.15 Terse output

For scripted usage where you typically want to generate tables or graphs of the results, fio can output the results in a semicolon separated format. The format is one long line of values, such as:

```
2;card0;0;0;7139336;121836;60004;1;10109;27.932460;116.933948;220;126861;3495.446807;
→1085.368601;226;126864;3523.635629;1089.012448;24063;99944;50.275485%;59818.274627;
→5540.657370;7155060;122104;60004;1;8338;29.086342;117.839068;388;128077;5032.488518;
→1234.785715;391;128085;5061.839412;1236.909129;23436;100928;50.287926%;59964.832030;
→5644.844189;14.595833%;19.394167%;123706;0;7313;0.1%;0.1%;0.1%;0.1%;0.1%;0.1%;100.0
→%;0.00%;0.00%;0.00%;0.00%;0.00%;0.00%;0.01%;0.02%;0.05%;0.16%;6.04%;40.40%;52.68%;0.
→64%;0.01%;0.00%;0.01%;0.00%;0.00%;0.00%;0.00%;0.00%
A description of this job goes here.
```

The job description (if provided) follows on a second line.

To enable terse output, use the `--minimal` or `--output-format`=terse command line options. The first value is the version of the terse output format. If the output has to be changed for some reason, this number will be incremented by 1 to signify that change.

Split up, the format is as follows (comments in brackets denote when a field was introduced or whether it's specific to some terse version):

```
terse version, fio version [v3], jobname, groupid, error
```

READ status:

```
Total IO (KiB), bandwidth (KiB/sec), IOPS, runtime (msec)
Submission latency: min, max, mean, stdev (usec)
Completion latency: min, max, mean, stdev (usec)
Completion latency percentiles: 20 fields (see below)
Total latency: min, max, mean, stdev (usec)
Bw (KiB/s): min, max, aggregate percentage of total, mean, stdev, number of␣
↪samples [v5]
IOPS [v5]: min, max, mean, stdev, number of samples
```

WRITE status:

```
Total IO (KiB), bandwidth (KiB/sec), IOPS, runtime (msec)
Submission latency: min, max, mean, stdev (usec)
Completion latency: min, max, mean, stdev (usec)
Completion latency percentiles: 20 fields (see below)
Total latency: min, max, mean, stdev (usec)
Bw (KiB/s): min, max, aggregate percentage of total, mean, stdev, number of␣
↪samples [v5]
IOPS [v5]: min, max, mean, stdev, number of samples
```

TRIM status [all but version 3]:

> Fields are similar to READ/WRITE status.

CPU usage:

```
user, system, context switches, major faults, minor faults
```

I/O depths:

```
<=1, 2, 4, 8, 16, 32, >=64
```

I/O latencies microseconds:

```
<=2, 4, 10, 20, 50, 100, 250, 500, 750, 1000
```

I/O latencies milliseconds:

```
<=2, 4, 10, 20, 50, 100, 250, 500, 750, 1000, 2000, >=2000
```

Disk utilization [v3]:

```
disk name, read ios, write ios, read merges, write merges, read ticks, write␣
↪ticks,
time spent in queue, disk utilization percentage
```

Additional Info (dependent on continue_on_error, default off):

```
total # errors, first error code
```

Additional Info (dependent on description being set):

```
Text description
```

Completion latency percentiles can be a grouping of up to 20 sets, so for the terse output fio writes all of them. Each field will look like this:

```
1.00%=6112
```

which is the Xth percentile, and the *usec* latency associated with it.

For *Disk utilization*, all disks used by fio are shown. So for each disk there will be a disk utilization section.

Below is a single line containing short names for each of the fields in the minimal output v3, separated by semicolons:

```
terse_version_3;fio_version;jobname;groupid;error;read_kb;read_bandwidth;read_iops;
→read_runtime_ms;read_slat_min;read_slat_max;read_slat_mean;read_slat_dev;read_clat_
→min;read_clat_max;read_clat_mean;read_clat_dev;read_clat_pct01;read_clat_pct02;read_
→clat_pct03;read_clat_pct04;read_clat_pct05;read_clat_pct06;read_clat_pct07;read_
→clat_pct08;read_clat_pct09;read_clat_pct10;read_clat_pct11;read_clat_pct12;read_
→clat_pct13;read_clat_pct14;read_clat_pct15;read_clat_pct16;read_clat_pct17;read_
→clat_pct18;read_clat_pct19;read_clat_pct20;read_tlat_min;read_lat_max;read_lat_mean;
→read_lat_dev;read_bw_min;read_bw_max;read_bw_agg_pct;read_bw_mean;read_bw_dev;write_
→kb;write_bandwidth;write_iops;write_runtime_ms;write_slat_min;write_slat_max;write_
→slat_mean;write_slat_dev;write_clat_min;write_clat_max;write_clat_mean;write_clat_
→dev;write_clat_pct01;write_clat_pct02;write_clat_pct03;write_clat_pct04;write_clat_
→pct05;write_clat_pct06;write_clat_pct07;write_clat_pct08;write_clat_pct09;write_
→clat_pct10;write_clat_pct11;write_clat_pct12;write_clat_pct13;write_clat_pct14;
→write_clat_pct15;write_clat_pct16;write_clat_pct17;write_clat_pct18;write_clat_
→pct19;write_clat_pct20;write_tlat_min;write_lat_max;write_lat_mean;write_lat_dev;
→write_bw_min;write_bw_max;write_bw_agg_pct;write_bw_mean;write_bw_dev;cpu_user;cpu_
→sys;cpu_csw;cpu_mjf;cpu_minf;iodepth_1;iodepth_2;iodepth_4;iodepth_8;iodepth_16;
→iodepth_32;iodepth_64;lat_2us;lat_4us;lat_10us;lat_20us;lat_50us;lat_100us;lat_
→250us;lat_500us;lat_750us;lat_1000us;lat_2ms;lat_4ms;lat_10ms;lat_20ms;lat_50ms;lat_
→100ms;lat_250ms;lat_500ms;lat_750ms;lat_1000ms;lat_2000ms;lat_over_2000ms;disk_name;
→disk_read_iops;disk_write_iops;disk_read_merges;disk_write_merges;disk_read_ticks;
→write_ticks;disk_queue_time;disk_util
```

1.16 JSON output

The *json* output format is intended to be both human readable and convenient for automated parsing. For the most part its sections mirror those of the *normal* output. The *runtime* value is reported in msec and the *bw* value is reported in 1024 bytes per second units.

1.17 JSON+ output

The *json+* output format is identical to the *json* output format except that it adds a full dump of the completion latency bins. Each *bins* object contains a set of (key, value) pairs where keys are latency durations and values count how many I/Os had completion latencies of the corresponding duration. For example, consider:

"bins" : { "87552" : 1, "89600" : 1, "94720" : 1, "96768" : 1, "97792" : 1, "99840" : 1, "100864" : 2, "103936" : 6, "104960" : 534, "105984" : 5995, "107008" : 7529, . . . }

This data indicates that one I/O required 87,552ns to complete, two I/Os required 100,864ns to complete, and 7529 I/Os required 107,008ns to complete.

Also included with fio is a Python script *fio_jsonplus_clat2csv* that takes json+ output and generates CSV-formatted latency data suitable for plotting.

The latency durations actually represent the midpoints of latency intervals. For details refer to `stat.h`.

1.18 Trace file format

There are two trace file format that you can encounter. The older (v1) format is unsupported since version 1.20-rc3 (March 2008). It will still be described below in case that you get an old trace and want to understand it.

In any case the trace is a simple text file with a single action per line.

1.18.1 Trace file format v1

Each line represents a single I/O action in the following format:

```
rw, offset, length
```

where *rw=0/1* for read/write, and the *offset* and *length* entries being in bytes.

This format is not supported in fio versions >= 1.20-rc3.

1.18.2 Trace file format v2

The second version of the trace file format was added in fio version 1.17. It allows to access more then one file per trace and has a bigger set of possible file actions.

The first line of the trace file has to be:

```
fio version 2 iolog
```

Following this can be lines in two different formats, which are described below.

The file management format:

```
filename action
```

The *filename* is given as an absolute path. The *action* can be one of these:

add Add the given *filename* to the trace.

open Open the file with the given *filename*. The *filename* has to have been added with the **add** action before.

close Close the file with the given *filename*. The file has to have been opened before.

The file I/O action format:

```
filename action offset length
```

The *filename* is given as an absolute path, and has to have been added and opened before it can be used with this format. The *offset* and *length* are given in bytes. The *action* can be one of these:

wait Wait for *offset* microseconds. Everything below 100 is discarded. The time is relative to the previous *wait* statement.

read Read *length* bytes beginning from *offset*.

write Write *length* bytes beginning from *offset*.

sync `fsync(2)` the file.

datasync `fdatasync(2)` the file.

trim Trim the given file from the given *offset* for *length* bytes.

1.19 I/O Replay - Merging Traces

Colocation is a common practice used to get the most out of a machine. Knowing which workloads play nicely with each other and which ones don't is a much harder task. While fio can replay workloads concurrently via multiple jobs, it leaves some variability up to the scheduler making results harder to reproduce. Merging is a way to make the order of events consistent.

Merging is integrated into I/O replay and done when a `merge_blktrace_file` is specified. The list of files passed to `read_iolog` go through the merge process and output a single file stored to the specified file. The output file is passed on as if it were the only file passed to `read_iolog`. An example would look like:

```
$ fio --read_iolog="<file1>:<file2>" --merge_blktrace_file="<output_file>"
```

Creating only the merged file can be done by passing the command line argument `merge-blktrace-only`.

Scaling traces can be done to see the relative impact of any particular trace being slowed down or sped up. `merge_blktrace_scalars` takes in a colon separated list of percentage scalars. It is index paired with the files passed to `read_iolog`.

With scaling, it may be desirable to match the running time of all traces. This can be done with `merge_blktrace_iters`. It is index paired with `read_iolog` just like `merge_blktrace_scalars`.

In an example, given two traces, A and B, each 60s long. If we want to see the impact of trace A issuing IOs twice as fast and repeat trace A over the runtime of trace B, the following can be done:

```
$ fio --read_iolog="<trace_a>":"<trace_b>" --merge_blktrace_file"<output_file>" --
→merge_blktrace_scalars="50:100" --merge_blktrace_iters="2:1"
```

This runs trace A at 2x the speed twice for approximately the same runtime as a single run of trace B.

1.20 CPU idleness profiling

In some cases, we want to understand CPU overhead in a test. For example, we test patches for the specific goodness of whether they reduce CPU usage. Fio implements a balloon approach to create a thread per CPU that runs at idle priority, meaning that it only runs when nobody else needs the cpu. By measuring the amount of work completed by the thread, idleness of each CPU can be derived accordingly.

An unit work is defined as touching a full page of unsigned characters. Mean and standard deviation of time to complete an unit work is reported in "unit work" section. Options can be chosen to report detailed percpu idleness or overall system idleness by aggregating percpu stats.

1.21 Verification and triggers

Fio is usually run in one of two ways, when data verification is done. The first is a normal write job of some sort with verify enabled. When the write phase has completed, fio switches to reads and verifies everything it wrote. The second model is running just the write phase, and then later on running the same job (but with reads instead of writes) to repeat the same I/O patterns and verify the contents. Both of these methods depend on the write phase being completed, as fio otherwise has no idea how much data was written.

With verification triggers, fio supports dumping the current write state to local files. Then a subsequent read verify workload can load this state and know exactly where to stop. This is useful for testing cases where power is cut to a server in a managed fashion, for instance.

A verification trigger consists of two things:

1) Storing the write state of each job.

2) Executing a trigger command.

The write state is relatively small, on the order of hundreds of bytes to single kilobytes. It contains information on the number of completions done, the last X completions, etc.

A trigger is invoked either through creation ('touch') of a specified file in the system, or through a timeout setting. If fio is run with `--trigger-file=` `/tmp/trigger-file`, then it will continually check for the existence of `/tmp/trigger-file`. When it sees this file, it will fire off the trigger (thus saving state, and executing the trigger command).

For client/server runs, there's both a local and remote trigger. If fio is running as a server backend, it will send the job states back to the client for safe storage, then execute the remote trigger, if specified. If a local trigger is specified, the server will still send back the write state, but the client will then execute the trigger.

1.21.1 Verification trigger example

Let's say we want to run a powercut test on the remote Linux machine 'server'. Our write workload is in `write-test.fio`. We want to cut power to 'server' at some point during the run, and we'll run this test from the safety or our local machine, 'localbox'. On the server, we'll start the fio backend normally:

```
server# fio --server
```

and on the client, we'll fire off the workload:

```
localbox$ fio --client=server --trigger-file=/tmp/my-trigger --trigger-remote="bash -
→c \"echo b > /proc/sysrq-triger\""
```

We set `/tmp/my-trigger` as the trigger file, and we tell fio to execute:

```
echo b > /proc/sysrq-trigger
```

on the server once it has received the trigger and sent us the write state. This will work, but it's not **really** cutting power to the server, it's merely abruptly rebooting it. If we have a remote way of cutting power to the server through IPMI or similar, we could do that through a local trigger command instead. Let's assume we have a script that does IPMI reboot of a given hostname, ipmi-reboot. On localbox, we could then have run fio with a local trigger instead:

```
localbox$ fio --client=server --trigger-file=/tmp/my-trigger --trigger="ipmi-reboot␣
→server"
```

For this case, fio would wait for the server to send us the write state, then execute `ipmi-reboot server` when that happened.

1.21.2 Loading verify state

To load stored write state, a read verification job file must contain the `verify_state_load` option. If that is set, fio will load the previously stored state. For a local fio run this is done by loading the files directly, and on a client/server run, the server backend will ask the client to send the files over and load them from there.

1.22 Log File Formats

Fio supports a variety of log file formats, for logging latencies, bandwidth, and IOPS. The logs share a common format, which looks like this:

> *time* (*msec*), *value*, *data direction*, *block size* (*bytes*), *offset* (*bytes*)

Time for the log entry is always in milliseconds. The *value* logged depends on the type of log, it will be one of the following:

Latency log Value is latency in nsecs

Bandwidth log Value is in KiB/sec

IOPS log Value is IOPS

Data direction is one of the following:

0 I/O is a READ

1 I/O is a WRITE

2 I/O is a TRIM

The entry's *block size* is always in bytes. The *offset* is the position in bytes from the start of the file for that particular I/O. The logging of the offset can be toggled with `log_offset`.

Fio defaults to logging every individual I/O but when windowed logging is set through `log_avg_msec`, either the average (by default) or the maximum (`log_max_value` is set) *value* seen over the specified period of time is recorded. Each *data direction* seen within the window period will aggregate its values in a separate row. Further, when using windowed logging the *block size* and *offset* entries will always contain 0.

1.23 Client/Server

Normally fio is invoked as a stand-alone application on the machine where the I/O workload should be generated. However, the backend and frontend of fio can be run separately i.e., the fio server can generate an I/O workload on the "Device Under Test" while being controlled by a client on another machine.

Start the server on the machine which has access to the storage DUT:

```
$ fio --server=args
```

where *args* defines what fio listens to. The arguments are of the form `type,hostname` or `IP,port`. *type* is either `ip` (or ip4) for TCP/IP v4, `ip6` for TCP/IP v6, or `sock` for a local unix domain socket. *hostname* is either a hostname or IP address, and *port* is the port to listen to (only valid for TCP/IP, not a local socket). Some examples:

1) `fio --server`

   Start a fio server, listening on all interfaces on the default port (8765).

2) `fio --server=ip:hostname,4444`

   Start a fio server, listening on IP belonging to hostname and on port 4444.

3) `fio --server=ip6:::1,4444`

   Start a fio server, listening on IPv6 localhost ::1 and on port 4444.

4) `fio --server=,4444`

   Start a fio server, listening on all interfaces on port 4444.

5) `fio --server=1.2.3.4`

   Start a fio server, listening on IP 1.2.3.4 on the default port.

6) `fio --server=sock:/tmp/fio.sock`

   Start a fio server, listening on the local socket `/tmp/fio.sock`.

Once a server is running, a "client" can connect to the fio server with:

```
fio <local-args> --client=<server> <remote-args> <job file(s)>
```

where *local-args* are arguments for the client where it is running, *server* is the connect string, and *remote-args* and *job file(s)* are sent to the server. The *server* string follows the same format as it does on the server side, to allow IP/hostname/socket and port strings.

Fio can connect to multiple servers this way:

```
fio --client=<server1> <job file(s)> --client=<server2> <job file(s)>
```

If the job file is located on the fio server, then you can tell the server to load a local file as well. This is done by using `--remote-config`

```
fio --client=server --remote-config /path/to/file.fio
```

Then fio will open this local (to the server) job file instead of being passed one from the client.

If you have many servers (example: 100 VMs/containers), you can input a pathname of a file containing host IPs/names as the parameter value for the `--client` option. For example, here is an example `host.list` file containing 2 hostnames:

```
host1.your.dns.domain
host2.your.dns.domain
```

The fio command would then be:

```
fio --client=host.list <job file(s)>
```

In this mode, you cannot input server-specific parameters or job files – all servers receive the same job file.

In order to let `fio --client` runs use a shared filesystem from multiple hosts, `fio --client` now prepends the IP address of the server to the filename. For example, if fio is using the directory `/mnt/nfs/fio` and is writing filename `fileio.tmp`, with a `--client` *hostfile* containing two hostnames h1 and h2 with IP addresses 192.168.10.120 and 192.168.10.121, then fio will create two files:

```
/mnt/nfs/fio/192.168.10.120.fileio.tmp
/mnt/nfs/fio/192.168.10.121.fileio.tmp
```

Examples

Some job file examples.

2.1 Poisson request flow

Download poisson-rate-submission.fio

```
[poisson-rate-submit]
size=128m
rw=randread
ioengine=libaio
iodepth=32
direct=1
# by setting the submit mode to offload, we can guarantee a fixed rate of
# submission regardless of what the device completion rate is.
io_submit_mode=offload
rate_iops=50
# Real world random request flow follows Poisson process. To give better
# insight on latency distribution, we simulate request flow under Poisson
# process.
rate_process=poisson
```

2.2 Latency profile

Download latency-profile.fio

```
# Test job that demonstrates how to use the latency target
# profiling. Fio will find the queue depth between 1..128
# that fits within the latency constraints of this 4k random
# read workload.
```

```
[global]
bs=4k
rw=randread
random_generator=lfsr
direct=1
ioengine=libaio
iodepth=128
# Set max acceptable latency to 500msec
latency_target=500000
# profile over a 5s window
latency_window=5000000
# 99.9% of IOs must be below the target
latency_percentile=99.9


[device]
filename=/dev/sda
```

2.3 Read 4 files with aio at different depths

Download aio-read.fio

```
; Read 4 files with aio at different depths
[global]
ioengine=libaio
buffered=0
rw=randread
bs=128k
size=512m
directory=/data1

[file1]
iodepth=4

[file2]
iodepth=32

[file3]
iodepth=8

[file4]
iodepth=16
```

2.4 Read backwards in a file

Download backwards-read.fio

```
# Demonstrates how to read backwards in a file.

[backwards-read]
bs=4k
# seek -8k back for every IO
rw=read:-8k
```

```
filename=128m
size=128m
```

2.5 Basic verification

Download `basic-verify.fio`

```
# The most basic form of data verification. Write the device randomly
# in 4K chunks, then read it back and verify the contents.
[write-and-verify]
rw=randwrite
bs=4k
direct=1
ioengine=libaio
iodepth=16
verify=crc32c
# Use /dev/XXX. For running this on a file instead, remove the filename
# option and add a size=32G (or whatever file size you want) instead.
filename=/dev/XXX
```

2.6 Fixed rate submission

Download `fixed-rate-submission.fio`

```
[fixed-rate-submit]
size=128m
rw=read
ioengine=libaio
iodepth=32
direct=1
# by setting the submit mode to offload, we can guarantee a fixed rate of
# submission regardless of what the device completion rate is.
io_submit_mode=offload
rate_iops=1000
```

2.7 Butterfly seek pattern

Download `butterfly.fio`

```
# Perform a butterfly/funnel seek pattern. This won't always alternate ends on
# every I/O but it will get close.

[global]
filename=/tmp/testfile
bs=4k
direct=1

[forward]
rw=read
```

```
flow=2
# Uncomment the size= and offset= lines to prevent each direction going past
# the middle of the file
#size=50%

[backward]
rw=read:-8k
flow=-2
#offset=50%
```

CHAPTER 3

TODO

3.1 GFIO TODO

In no particular order:

- Ability to save job files. Probably in an extended gfio format, so we can include options/settings outside of a fio job file.

- End view improvements:

  - Cleanup the layout

  - Add ability to save the results

  - Add ability to load end-results as well

  - Add ability to request graphs of whatever graphing options the fio job included.

  - Add ability to graph completion latencies, percentiles, etc.

- Add ability to edit job options:

  - We need an options view after sending a job, that allows us to visually see what was parsed, make changes, resubmit.

  - Job options are already converted across the network and are available in gfio_client->o for view/edit. We'll need a FIO_NET_CMD_UPDATE_OPTIONS command to send them back, and backend support for updating an existing set of options.

- Add support for printing end results, graphs, etc.

- Improve the auto-start backend functionality, it's quite buggy.

- Ensure that it works on OSX and Windows. We'll need a bit of porting work there.

- Persistent store of prefences set. This will need a per-OS bit as well, using gfonf on Linux, registry on Windows, ?? on OSX.

- Ensure that local errors go to our log, instead of being displayed on the console.

- Ensure that the whole connect/send/start button logic is sane. Right now it works when you perform the right sequence, but if you connect and disconnect, things can get confused. We'll need to improve how we store and send job files. Right now they are in ge->job_files[] and are always emptied on send. Keep them around?

- Commit rate display is not enabled.

- Group status reporting is not enabled.

- Split gfio.c a bit. Add gfio/ sub directory, and split it into files based on functionality. It's already ~3000 lines long.

- Attempt to ensure that we work with gtk 2.10 and newer. Right now the required version is ~2.18 (not quite known).

3.2 Server TODO

- Collate ETA output from multiple connections into 1
- If group_reporting is set, collate final output from multiple connections

3.3 Steady State TODO

Known issues/TODO (for steady-state)

- Allow user to specify the frequency of measurements
- Better documentation for output
- Report read, write, trim IOPS/BW separately
- Semantics for the ring buffer ss->head are confusing. ss->head points to the beginning of the buffer up through the point where the buffer is filled for the first time. afterwards, when a new element is added, ss->head is advanced to point to the second element in the buffer. if steady state is attained upon adding a new element, ss->head is not advanced so it actually does point to the head of the buffer.

CHAPTER 4

Moral License

As specified by the COPYING file, fio is free software published under version 2 of the GPL license. That covers the copying part of the license. When using fio, you are encouraged to uphold the following moral obligations:

- If you publish results that are done using fio, it should be clearly stated that fio was used. The specific version should also be listed.

- If you develop features or bug fixes for fio, they should be sent upstream for inclusion into the main repository. This isn't specific to fio, that is a general rule for any open source project. It's just the Right Thing to do. Plus it means that you don't have to maintain the feature or change internally. In the long run, this is saving you a lot of time.

I would consider the above to fall under "common courtesy", but since people tend to have differing opinions of that, it doesn't hurt to spell out my expectations clearly.

License

```
                    GNU GENERAL PUBLIC LICENSE
                       Version 2, June 1991

 Copyright (C) 1989, 1991 Free Software Foundation, Inc.,
 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA
 Everyone is permitted to copy and distribute verbatim copies
 of this license document, but changing it is not allowed.

                          Preamble

  The licenses for most software are designed to take away your
freedom to share and change it.  By contrast, the GNU General Public
License is intended to guarantee your freedom to share and change free
software--to make sure the software is free for all its users.  This
General Public License applies to most of the Free Software
Foundation's software and to any other program whose authors commit to
using it.  (Some other Free Software Foundation software is covered by
the GNU Lesser General Public License instead.)  You can apply it to
your programs, too.

  When we speak of free software, we are referring to freedom, not
price.  Our General Public Licenses are designed to make sure that you
have the freedom to distribute copies of free software (and charge for
this service if you wish), that you receive source code or can get it
if you want it, that you can change the software or use pieces of it
in new free programs; and that you know you can do these things.

  To protect your rights, we need to make restrictions that forbid
anyone to deny you these rights or to ask you to surrender the rights.
These restrictions translate to certain responsibilities for you if you
distribute copies of the software, or if you modify it.

  For example, if you distribute copies of such a program, whether
gratis or for a fee, you must give the recipients all the rights that
```

```
you have.  You must make sure that they, too, receive or can get the
source code.  And you must show them these terms so they know their
rights.

  We protect your rights with two steps: (1) copyright the software, and
(2) offer you this license which gives you legal permission to copy,
distribute and/or modify the software.

  Also, for each author's protection and ours, we want to make certain
that everyone understands that there is no warranty for this free
software.  If the software is modified by someone else and passed on, we
want its recipients to know that what they have is not the original, so
that any problems introduced by others will not reflect on the original
authors' reputations.

  Finally, any free program is threatened constantly by software
patents.  We wish to avoid the danger that redistributors of a free
program will individually obtain patent licenses, in effect making the
program proprietary.  To prevent this, we have made it clear that any
patent must be licensed for everyone's free use or not licensed at all.

  The precise terms and conditions for copying, distribution and
modification follow.

                    GNU GENERAL PUBLIC LICENSE
   TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

  0. This License applies to any program or other work which contains
a notice placed by the copyright holder saying it may be distributed
under the terms of this General Public License.  The "Program", below,
refers to any such program or work, and a "work based on the Program"
means either the Program or any derivative work under copyright law:
that is to say, a work containing the Program or a portion of it,
either verbatim or with modifications and/or translated into another
language.  (Hereinafter, translation is included without limitation in
the term "modification".)  Each licensee is addressed as "you".

Activities other than copying, distribution and modification are not
covered by this License; they are outside its scope.  The act of
running the Program is not restricted, and the output from the Program
is covered only if its contents constitute a work based on the
Program (independent of having been made by running the Program).
Whether that is true depends on what the Program does.

  1. You may copy and distribute verbatim copies of the Program's
source code as you receive it, in any medium, provided that you
conspicuously and appropriately publish on each copy an appropriate
copyright notice and disclaimer of warranty; keep intact all the
notices that refer to this License and to the absence of any warranty;
and give any other recipients of the Program a copy of this License
along with the Program.

You may charge a fee for the physical act of transferring a copy, and
you may at your option offer warranty protection in exchange for a fee.

  2. You may modify your copy or copies of the Program or any portion
of it, thus forming a work based on the Program, and copy and
```

```
distribute such modifications or work under the terms of Section 1
above, provided that you also meet all of these conditions:

    a) You must cause the modified files to carry prominent notices
    stating that you changed the files and the date of any change.

    b) You must cause any work that you distribute or publish, that in
    whole or in part contains or is derived from the Program or any
    part thereof, to be licensed as a whole at no charge to all third
    parties under the terms of this License.

    c) If the modified program normally reads commands interactively
    when run, you must cause it, when started running for such
    interactive use in the most ordinary way, to print or display an
    announcement including an appropriate copyright notice and a
    notice that there is no warranty (or else, saying that you provide
    a warranty) and that users may redistribute the program under
    these conditions, and telling the user how to view a copy of this
    License.  (Exception: if the Program itself is interactive but
    does not normally print such an announcement, your work based on
    the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole.  If
identifiable sections of that work are not derived from the Program,
and can be reasonably considered independent and separate works in
themselves, then this License, and its terms, do not apply to those
sections when you distribute them as separate works.  But when you
distribute the same sections as part of a whole which is a work based
on the Program, the distribution of the whole must be on the terms of
this License, whose permissions for other licensees extend to the
entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest
your rights to work written entirely by you; rather, the intent is to
exercise the right to control the distribution of derivative or
collective works based on the Program.

In addition, mere aggregation of another work not based on the Program
with the Program (or with a work based on the Program) on a volume of
a storage or distribution medium does not bring the other work under
the scope of this License.

  3. You may copy and distribute the Program (or a work based on it,
under Section 2) in object code or executable form under the terms of
Sections 1 and 2 above provided that you also do one of the following:

    a) Accompany it with the complete corresponding machine-readable
    source code, which must be distributed under the terms of Sections
    1 and 2 above on a medium customarily used for software interchange; or,

    b) Accompany it with a written offer, valid for at least three
    years, to give any third party, for a charge no more than your
    cost of physically performing source distribution, a complete
    machine-readable copy of the corresponding source code, to be
    distributed under the terms of Sections 1 and 2 above on a medium
    customarily used for software interchange; or,
```

```
      c) Accompany it with the information you received as to the offer
      to distribute corresponding source code.  (This alternative is
      allowed only for noncommercial distribution and only if you
      received the program in object code or executable form with such
      an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for
making modifications to it.  For an executable work, complete source
code means all the source code for all modules it contains, plus any
associated interface definition files, plus the scripts used to
control compilation and installation of the executable.  However, as a
special exception, the source code distributed need not include
anything that is normally distributed (in either source or binary
form) with the major components (compiler, kernel, and so on) of the
operating system on which the executable runs, unless that component
itself accompanies the executable.

If distribution of executable or object code is made by offering
access to copy from a designated place, then offering equivalent
access to copy the source code from the same place counts as
distribution of the source code, even though third parties are not
compelled to copy the source along with the object code.

  4. You may not copy, modify, sublicense, or distribute the Program
except as expressly provided under this License.  Any attempt
otherwise to copy, modify, sublicense or distribute the Program is
void, and will automatically terminate your rights under this License.
However, parties who have received copies, or rights, from you under
this License will not have their licenses terminated so long as such
parties remain in full compliance.

  5. You are not required to accept this License, since you have not
signed it.  However, nothing else grants you permission to modify or
distribute the Program or its derivative works.  These actions are
prohibited by law if you do not accept this License.  Therefore, by
modifying or distributing the Program (or any work based on the
Program), you indicate your acceptance of this License to do so, and
all its terms and conditions for copying, distributing or modifying
the Program or works based on it.

  6. Each time you redistribute the Program (or any work based on the
Program), the recipient automatically receives a license from the
original licensor to copy, distribute or modify the Program subject to
these terms and conditions.  You may not impose any further
restrictions on the recipients' exercise of the rights granted herein.
You are not responsible for enforcing compliance by third parties to
this License.

  7. If, as a consequence of a court judgment or allegation of patent
infringement or for any other reason (not limited to patent issues),
conditions are imposed on you (whether by court order, agreement or
otherwise) that contradict the conditions of this License, they do not
excuse you from the conditions of this License.  If you cannot
distribute so as to satisfy simultaneously your obligations under this
License and any other pertinent obligations, then as a consequence you
may not distribute the Program at all.  For example, if a patent
license would not permit royalty-free redistribution of the Program by
```

```
all those who receive copies directly or indirectly through you, then
the only way you could satisfy both it and this License would be to
refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under
any particular circumstance, the balance of the section is intended to
apply and the section as a whole is intended to apply in other
circumstances.

It is not the purpose of this section to induce you to infringe any
patents or other property right claims or to contest validity of any
such claims; this section has the sole purpose of protecting the
integrity of the free software distribution system, which is
implemented by public license practices.  Many people have made
generous contributions to the wide range of software distributed
through that system in reliance on consistent application of that
system; it is up to the author/donor to decide if he or she is willing
to distribute software through any other system and a licensee cannot
impose that choice.

This section is intended to make thoroughly clear what is believed to
be a consequence of the rest of this License.

  8. If the distribution and/or use of the Program is restricted in
certain countries either by patents or by copyrighted interfaces, the
original copyright holder who places the Program under this License
may add an explicit geographical distribution limitation excluding
those countries, so that distribution is permitted only in or among
countries not thus excluded.  In such case, this License incorporates
the limitation as if written in the body of this License.

  9. The Free Software Foundation may publish revised and/or new versions
of the General Public License from time to time.  Such new versions will
be similar in spirit to the present version, but may differ in detail to
address new problems or concerns.

Each version is given a distinguishing version number.  If the Program
specifies a version number of this License which applies to it and "any
later version", you have the option of following the terms and conditions
either of that version or of any later version published by the Free
Software Foundation.  If the Program does not specify a version number of
this License, you may choose any version ever published by the Free Software
Foundation.

  10. If you wish to incorporate parts of the Program into other free
programs whose distribution conditions are different, write to the author
to ask for permission.  For software which is copyrighted by the Free
Software Foundation, write to the Free Software Foundation; we sometimes
make exceptions for this.  Our decision will be guided by the two goals
of preserving the free status of all derivatives of our free software and
of promoting the sharing and reuse of software generally.

                          NO WARRANTY

  11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY
FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW.  EXCEPT WHEN
OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES
```

```
PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED
OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF
MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE.  THE ENTIRE RISK AS
TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU.  SHOULD THE
PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING,
REPAIR OR CORRECTION.

  12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING
WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR
REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES,
INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING
OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED
TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY
YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER
PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE
POSSIBILITY OF SUCH DAMAGES.

                    END OF TERMS AND CONDITIONS

          How to Apply These Terms to Your New Programs

  If you develop a new program, and you want it to be of the greatest
possible use to the public, the best way to achieve this is to make it
free software which everyone can redistribute and change under these terms.

  To do so, attach the following notices to the program.  It is safest
to attach them to the start of each source file to most effectively
convey the exclusion of warranty; and each file should have at least
the "copyright" line and a pointer to where the full notice is found.

    <one line to give the program's name and a brief idea of what it does.>
    Copyright (C) <year>  <name of author>

    This program is free software; you can redistribute it and/or modify
    it under the terms of the GNU General Public License as published by
    the Free Software Foundation; either version 2 of the License, or
    (at your option) any later version.

    This program is distributed in the hope that it will be useful,
    but WITHOUT ANY WARRANTY; without even the implied warranty of
    MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
    GNU General Public License for more details.

    You should have received a copy of the GNU General Public License along
    with this program; if not, write to the Free Software Foundation, Inc.,
    51 Franklin Street, Fifth Floor, Boston, MA 02110-1301 USA.

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this
when it starts in an interactive mode:

    Gnomovision version 69, Copyright (C) year name of author
    Gnomovision comes with ABSOLUTELY NO WARRANTY; for details type `show w'.
    This is free software, and you are welcome to redistribute it
    under certain conditions; type `show c' for details.
```

```
The hypothetical commands `show w' and `show c' should show the appropriate
parts of the General Public License.  Of course, the commands you use may
be called something other than `show w' and `show c'; they could even be
mouse-clicks or menu items--whatever suits your program.

You should also get your employer (if you work as a programmer) or your
school, if any, to sign a "copyright disclaimer" for the program, if
necessary.  Here is a sample; alter the names:

  Yoyodyne, Inc., hereby disclaims all copyright interest in the program
  `Gnomovision' (which makes passes at compilers) written by James Hacker.

  <signature of Ty Coon>, 1 April 1989
  Ty Coon, President of Vice

This General Public License does not permit incorporating your program into
proprietary programs.  If your program is a subroutine library, you may
consider it more useful to permit linking proprietary applications with the
library.  If this is what you want to do, use the GNU Lesser General
Public License instead of this License.
```

Indices and tables

- genindex
- search

L

lat_percentiles=bool
latency_percentile=float
latency_target=time
latency_window=time
listen : [netsplice] [net]
lockfile=str
lockmem=int
log_avg_msec=int
log_compression=int
log_compression_cpus=str
log_hist_coarseness=int
log_hist_msec=int
log_max_value=bool
log_offset=bool
log_store_compressed=bool
log_unix_epoch=bool
loops=int

M

max_latency=time
max_open_zones=int
merge_blktrace_file=str
merge_blktrace_iters=float_list
merge_blktrace_scalars=float_list
mss : [netsplice] [net]

N

name=str
namenode=str :   [libhdfs]

new_group
nice=int
nodelay=bool : [netsplice] [net]
norandommap
nrfiles=int
numa_cpu_nodes=str
numa_mem_policy=str
number_ios=int
numjobs=int

O

offset=int
offset_align=int
offset_increment=int
opendir=str
openfiles=int
overwrite=bool

P

per_job_logs=bool
percentage_random=int[,int][,int]
percentile_list=float_list
pingpong : [netsplice] [net]
pool=str : [rbd,rados]
port=int
pre_read=bool
prio=int
prioclass=int
profile=str